

Computer Architecture (Practical Class)

Assembly: Arithmetic Operations

Luís Nogueira

Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto

lmn@isep.ipp.pt

2023/2024

- Most of the operations are given as instruction classes, as they can have different variants with different operand sizes
 - For example, the instruction class add consists of four addition instructions: addb, addw, addl, and addq, adding bytes, words, double words, and quad words, respectively
- Recall that type conversion has higher precedence than arithmetic operations
 - We can only operate on arguments of the same size
 - Moving from a smaller to a larger data size can involve either sign extension (for signed values) or zero extension (for unsigned values)
- Do not forget that instructions that move or generate 32-bit register values also set the upper 32 bits of the register to zero

- The ADD instruction adds two integers
- Usage: `add origin, destination`
- Performs the operation $destination = destination + origin$ (the result is placed in *destination*)
- *origin* can be a memory address, a constant value or a register
- *destination* can be a memory address or a register
- As for any other instruction, a memory address for *origin* and *destination* cannot be used simultaneously
- The ADD instruction can add numbers of 8(b), 16(w), 32(l), or 64(q) bits

Adding bytes, words, double words (longs), and quads

```
addb $10, %al          # adds 10 to the 8-bit AL register; AL=AL+10
addw %bx, %cx          # adds the value of BX to CX (16 bits); CX=CX+BX
addl var1(%rip), %eax  # adds the 32-bit value in var1 to EAX; EAX=EAX+var1
addl %eax, %eax        # adds EAX to itself; EAX=EAX+EAX
addq %rcx, %rax        # adds RCX to RAX; RAX=RAX+RCX
```

- The SUB instruction subtracts two integers
- Usage: `sub origin, destination`
- Performs the operation $destination = destination - origin$ (the result is placed in $destination$)
- $origin$ can be a memory address, a constant value or a register
- $destination$ can be a memory address or a register
- A memory address for $origin$ and $destination$ cannot be used simultaneously
- The SUB instruction can subtract numbers of 8(b), 16(w), 32(l), or 64(q) bits

Subtracting bytes, words, double words (longs), and quads

```
subl $10, %eax      # subtract 10 to the current value of EAX; EAX=EAX-10
subw %bx, %cx      # subtract the value of BX to CX (16 bits); CX=CX-BX
subb var1(%rip), %al  # subtract the 8-bit value in var1 to AL; AL=AL-var1
subq %rcx, %rax    # subtract RCX to RAX; RAX=RAX-RCX
```

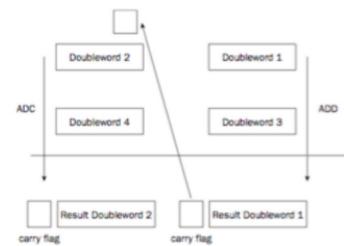
- The INC and DEC instructions increment (INC) and decrement (DEC) an integer by one, respectively
- Usage: `inc destination`
`dec destination`
- performs the operation $destination = destination +/- 1$;
- $destination$ can be a memory address or a register
- the INC and DEC instructions can be used in numbers of 8(b), 16(w), 32(l) or 64(q) bits

Increment/Decrement bytes, words, double words (longs), and quads

```
incq %rax    # RAX=RAX+1 (64 bits)
incl %eax    # EAX=EAX+1 (32 bits)
incw %bx     # BX=BX+1 (16 bits)
decb %cl     # CL=CL-1 (8 bits)
```

adc *origin*, *destination*

- The ADC instruction can be used to add two integer values, along with the value contained in the *carry flag* set by a previous addition
- If your program treats the bits in a register as *unsigned* numbers, you must watch to see if your arithmetic sets the carry flag on, indicating the result is wrong (more on this in the next class)
- Performs the operation: $\text{destination} = \text{destination} + \text{origin} + CF$
- The ADC instruction can add numbers of 8(b), 16(w), 32(l), or 64(q) bits



Add With Carry (ADC) Example

```
.global adctest
adctest:
...
    movb $0xFF, %al
    movb $0x1, %ah
    movb $0x1, %cl
    movb $0x0, %ch

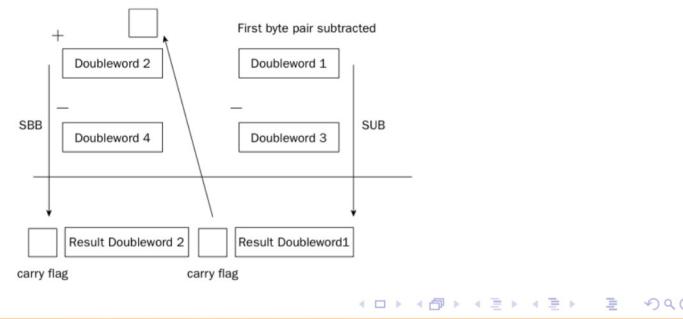
    # cl = cl + al (8 bits)
    addb %al, %cl

    # ch = ch + ah + CF (8 bits)
    adcb %ah, %ch

...
ret
```

sbb *origin*, *destination*

- The SBB instruction can be used to subtract two integer values, along with the value contained in the *carry flag* set by a previous subtraction
- If your program treats the bits in a word as *unsigned* numbers, you must watch to see if your arithmetic sets the carry flag on, indicating the result is wrong (more on this in the next class)
- Performs the operation: $destination = destination - (origin + CF)$
- The SBB instruction can subtract numbers of 8(b), 16(w), 32(l), or 64(q) bits.



- When mixing data of different sizes, you must make the right choice of arithmetic instructions, sign extensions, and zero extensions. These depend on subtle aspects of type conversion and the interactions between the arithmetic instructions
- This is illustrated by the following C function:

Mixing different sizes in arithmetic instructions

```
int x,y; // 32 bits (signed)  
  
long fun(){  
    long t1 = (long)x + y;  
    long t2 = (long)(x+y);  
    return t1 - t2;  
}
```

- The two additions in this function proceed as follows:
 - Recall that type conversion has higher precedence than addition, and so the first addition calls for x to be converted to 64 bits, and by operand promotion y is also converted. Value $t1$ is then computed using 64-bit addition
 - On the other hand, $t2$ is computed by performing a 32-bit addition and then extending this value to 64 bits

- The equivalent Assembly code for this function is as follows:

Mixing different sizes in arithmetic instructions

```
fun:  
    movslq x(%rip),%rax      # Convert x to long  
    movslq y(%rip),%rdx      # Convert y to long  
    addq %rdx,%rax          # t1 (64-bit addition)  
    movl y(%rip),%edx        # Move y to %edx (lower 32 bits of RDX)  
    addl x(%rip),%edx        # t2 (32-bit addition)  
    movslq %edx,%rdx         # Convert t2 to long  
    subq %rdx,%rax           # Return t1 - t2  
    ret
```

- The two additions in this function proceed as follows:

- t_1 is computed by first sign extending the arguments. The `movslq` instructions take the 32 bits of x and y and sign extend them to 64 bits in registers `%rax` and `%rdx`. The `addq` instruction then performs 64-bit addition to get t_1
- t_2 is computed by performing 32-bit addition with x and y . This value is sign extended to 64 bits (within a single register) to get t_2

- Multiplying two 64-bit *signed* or *unsigned* integers can yield a product that requires 128 bits to represent. Intel refers to a 16-byte quantity as an *oct word*
- There are “one-operand” instructions that support generating the full 128-bit product of two 64-bit numbers
 - One argument must be in register %rax, and the other is given as the instruction source operand
 - The product is then stored in registers %rdx (high-order 64 bits) and %rax (low-order 64 bits)
- There is also a “two-operand” multiply instruction, generating a 64-bit product from two 64-bit operands
 - Recall that when truncating the product to 64 bits, both unsigned multiply and two’s-complement multiply have the same bit-level behaviour

mul *origin*

- The MUL instruction multiplies two **unsigned** integers
- Performs the operation: $destination = origin \times operand2$
- *origin* can be a memory address or a register
- *operand2* is the RAX, EAX, AX or AL register, depending on the size of *origin*
- *destination* is the RDX:RAX, EDX:EAX, DX:AX or AX registers, depending on the size of *origin*;

Size of <i>origin</i>	<i>operand2</i>	<i>destination</i>
8 bits	AL	AX
16 bits	AX	DX:AX
32 bits	EAX	EDX:EAX
64 bits	RAX	RDX:RAX

- The MUL instruction can multiply numbers of 8(b), 16(w), 32(l), or 64(q) bits.

Unsigned Multiplication Example

```
.global multest
multest:
...
    movw $200, %ax
    movw $2, %cx
    # multiply %cx by %ax
    # result in %dx:%ax
    mulw %cx
...
ret
```

`imul origin`
`imul origin, destination`

- The IMUL instruction multiplies two **signed** integers
- `imul origin`
 - The use of registers is the same as presented in MUL
- `imul origin, destination`
 - Performs $destination = destination \times origin$
 - *origin* can be a memory address, a constant value or a register (16, 32, or 64-bit)
 - *destination* can be a 16,32, or 64-bit register

Signed Multiplication Example

```
.global imultest
imultest:
# ...
    movl $200, %eax
    movl $-2, %ecx
    imull %ecx          # %edx:%eax = %ecx * %eax

    movq $-123, %rcx
    imulq $4, %rcx      # %rcx = 4 * %rcx
# ...
    ret
```

- Signed and unsigned divisions are only supported through single-operand instructions
- Their behaviour is similar to the single-operand multiply instructions
 - The dividend is the 128-bit value stored in registers `%rdx` (high-order 64 bits) and `%rax` (low-order 64 bits)
 - The divisor is given as the instruction operand
 - The instruction stores the quotient in register `%rax` and the remainder in register `%rdx`
- For most practical cases, the dividend is given as a value up to 64 bits
 - This value should be stored in register `%rax`
 - The bits of `%rdx` should then be set to either all zeros (unsigned arithmetic) or the sign bit of `%rax` (signed arithmetic)

div divisor
idiv divisor

- The DIV/IDIV instruction is used for unsigned-signed division
- Performs the operations:
 - $\text{quotient} = \text{dividend} \div \text{divisor}$
 - $\text{remainder} = \text{dividend} \bmod \text{divisor}$
- *divisor* can be a memory address or a register
- *dividend* is the RDX: RAX, EDX:EAX, DX:AX or AX register, depending on the size of *divisor*
- The *quotient* and *remainder* of the division are put in different sections of the RAX and RDX registers, depending on the size of *divisor*

Size of divisor	dividend	quotient	remainder
8 bits	AX	AL	AH
16 bits	DX:AX	AX	DX
32 bits	EDX:EAX	EAX	EDX
64 bits	RDX:RAX	RAX	RDX

Unsigned Division Example

```
.global divtest
divtest:
...
# dividend: ax
movw $100, %ax
# divisor: cl
movb $3, %cl
# divides %ax by %cl
# remainder in %ah
# quotient in %al
divb %cl
...
ret
```

Important note

- Always initialize **all** bits of the dividend!

Very bad idea!

```
.global bad_div

bad_div:
...
# dividend: eax
movl $100, %eax
# divisor: ecx
movl $3, %ecx

# divides %edx:%eax by %ecx
# Problem: the unknown content in %edx becomes part of the dividend

# remainder in %edx
# quotient in %eax
divl %ecx

...
ret
```



- Preparing the dividend depends on whether unsigned (*div*) or signed (*idiv*) division is to be performed
- In the former case, the higher order part of the dividend must be set to 0
- In the latter case, a set of instructions is used to produce a correct dividend before a division instruction through sign extension
- The Intel-syntax conversion instructions:
 - *cbw* — sign-extend byte in *%al* to word in *%ax*
 - *cwde* — sign-extend word in *%ax* to double word (long) in *%eax*
 - *cwd* — sign-extend word in *%ax* to double word (long) in *%dx : %ax*
 - *cdq* — sign-extend double word (long) in *%eax* to quad word in *%edx : %eax*
 - *cdqe* — sign-extend double word (long) in *%eax* to quad word in *%rax*
 - *cqo* — sign-extend quad in *%rax* to oct word in *%rdx : %rax*
- are called *cbtw*, *cwtl*, *cwtd*, *cltd*, *cltq*, and *cqto* in AT&T naming
 - This is one of the few instructions whose GAS name is very different from the Intel version. GAS accepts either mnemonic, but Intel-syntax assemblers like NASM may only accept the Intel names

Correct sign extension

```
.global div_ok

div_ok:
...
# dividend: %eax
movl $-100, %eax
# converts the signed long in %eax to the signed double long in %edx:%eax
cltd

# divisor: ecx
movl $3, %ecx

# divides %edx:%eax by %ecx (remainder in %edx, quotient in %eax)
idivl %ecx
...
ret
```

A C function *arithprob* that accesses the global variables *a*, *b*, *c*, and *d* has the following body:

```
return a*b + c*d;
```

It compiles to the following Assembly code:

arithprob.s

```
arithprob:  
    movl    d(%rip), %ecx  
    movslq %ecx,%rcx  
    movb    b(%rip),%sil  
    movsb1 %sil,%esi  
    imulq  c(%rip),%rcx  
    imull  a(%rip),%esi  
    movslq %esi,%rsi  
    addq    %rsi,%rcx  
    movq    %rcx,%rax  
    ret
```

Based on this Assembly code, determine the types of each of the global variables and write a function prototype describing the return type for *arithprob*. The return of an Assembly function up to a 64-bit value is placed in *%rax* (or parts of it)