

Computer Architecture (Practical Class)

Assembly: Stack and Functions - Parameters and Local Variables

Luís Nogueira

Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto

lmn@isep.ipp.pt

2023/2024

- Many (if not most) functions require some sort of input data
- In x86-64, most of the data passed to procedures take place via registers
- Up to six integral (i.e., integer and pointer) arguments can be passed via registers
- When a function has more than six integral arguments, the other ones are passed on the stack
- This method works equally as well for any Assembly program, even if it wasn't derived from a C program

- The registers are used in a **specified order**, with the name used for a register depending on the size of the data type being passed

| Operand size (bits) | Argument number | | | | | |
|------------------------|-----------------|------|------|------|------|------|
| | 1 | 2 | 3 | 4 | 5 | 6 |
| 64 | %rdi | %rsi | %rdx | %rcx | %r8 | %r9 |
| 32 | %edi | %esi | %edx | %ecx | %r8d | %r9d |
| 16 | %di | %si | %dx | %cx | %r8w | %r9w |
| 8 | %dil | %sil | %dl | %cl | %r8b | %r9b |

- Consider the following functions:

Function f1

```
long f1(long a, long b, long c){  
    return a + b + f2(c);  
}
```

Function f2

```
long f2(long c){  
    return c + 1;  
}
```

- How can you implement an equivalent code in Assembly?

Function f1

```
long f1(long a, long b, long c){  
    return a + b + f2(c);  
}
```

Function f2

```
long f2(long c){  
    return c + 1;  
}
```

Function f1

```
f1:  
# a in %rdi, b in %rsi, c in %rdx  
# a + b  
addq %rsi, %rdi  
# store %rdi on stack  
pushq %rdi  
# prepare function call f2(c)  
movq %rdx, %rdi  
call f2  
# restore %rdi  
popq %rdi  
# a + b + f2(c)  
addq %rdi, %rax  
ret
```

Function f2

```
f2:  
# c in %rdi  
# c + 1  
incq %rdi  
movq %rdi, %rax  
ret
```

- Consider the following functions:

Function f1

```
void f1(int a, char b, int *res){  
    *res = f2(a,b);  
}
```

Function f2

```
int f2(int a, char b){  
    return a * b;  
}
```

- Implement an equivalent code in Assembly

Function f1

```
void f1(int a, char b, int *res){  
    *res = f2(a,b);  
}
```

Function f2

```
int f2(int a, char b){  
    return a * b;  
}
```

Function f1

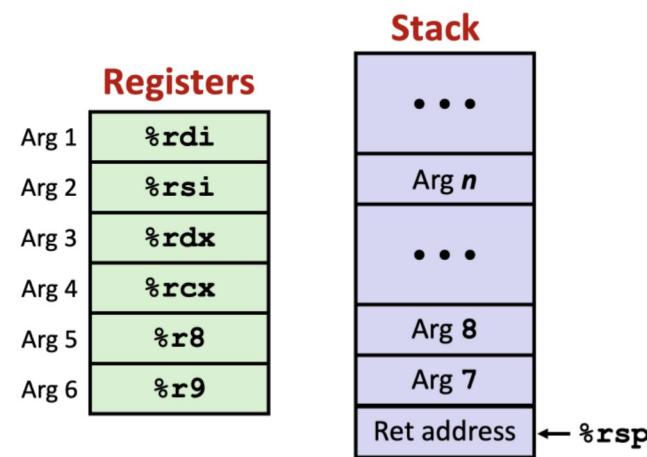
```
f1:  
# a in %edi, b in %sil  
# *res in %rdx  
# store %rdx on stack  
pushq %rdx  
# f2(a,b)  
call f2  
#restore %rdx  
popq %rdx  
movl %eax, (%rdx)  
ret
```

Function f2

```
f2:  
# a in %edi, b in %sil  
# sign extend b  
movsbl %sil, %eax  
  
# a * b  
# result in %edx:%eax  
imull %edi  
ret
```

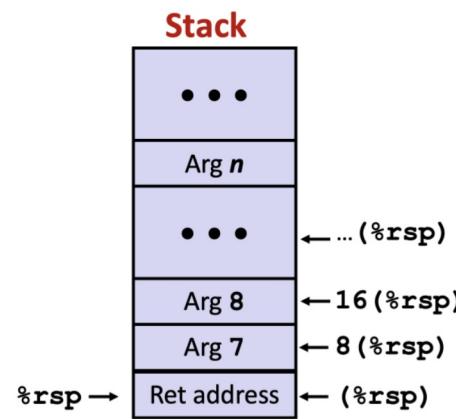
- Assume that procedure P calls procedure Q with n integral arguments, such that $n > 6$
- P , the caller, must copy arguments 1 to 6 into the appropriate registers, and put arguments 7 through n onto the stack, with **argument 7 at the top of the stack**
- When passing parameters on the stack, all data sizes are rounded up to be multiples of eight
- With the arguments in place, procedure P can then execute a *call* instruction to transfer control to procedure Q
 - When the *call* instruction is executed, it places the return address from P onto the top of the stack as well, so Q knows where to return
- Q , the callee, can access its first six arguments via registers and the remaining ones from the stack

- All these details leave the registers and the stack in the state shown here



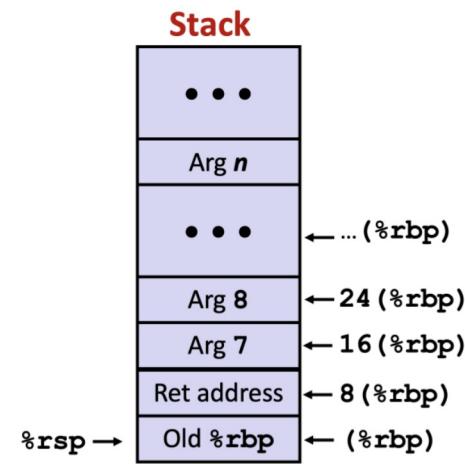
- The stack pointer (RSP) points to the top of the stack, where the return address is located
 - All of the input parameters for the function are located “underneath” the return address on the stack
- Popping values off of the stack to retrieve the input parameters would cause a problem, as the return address might be lost in the process
- Each parameter can be indirectly accessed via the offset from the RSP register, without having to POP values off of the stack
 - `movq 8(%rsp), %rax` - copies (without popping!) the 7th parameter to the RAX register

- Using indirect addressing with the RSP register to access the 7+ input parameters



- There is an inconvenient problem with this approach! What happens if data is pushed to the stack throughout the callee's code?
- If this happens, it would change the location of the RSP stack pointer and throw off the indirect addressing values for accessing the parameters in the stack
- To avoid this problem, it is common practice to copy the RSP register value to the RBP register when entering the function
- This ensures that there is a register that always contains the correct pointer to the top of the stack when the function is called (a push does not changes RBP)
- To avoid corrupting the RBP register on the caller, the callee (before the RSP register value is copied) should also place the RBP register value on the stack (RBP is a callee-saved register) and restore it before returning

- All these details leave the stack in the state shown here



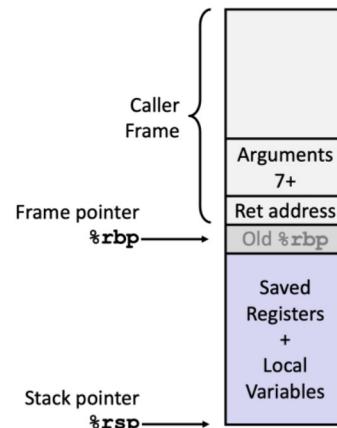
- This approach has created a standard set of instructions that are found in procedures that make a heavy use of the stack
- Stack space is allocated in **frames**
 - A section of stack used by one procedure call to store context while running, delimited by `%rbp` and `%rsp`
- Procedure code manages stack frames explicitly
 - **Prologue/Setup:** allocate space at start of procedure
 - **Epilogue/Cleanup:** deallocate space before return

- The base pointer `%rbp` being a stable “anchor” to the beginning of the stack frame throughout the execution of a function, is very convenient for manual Assembly coding and for debugging

Function template with Prologue and Epilogue

```
function:  
    # prologue  
    pushq %rbp      # save the original value of RBP  
    movq %rsp, %rbp # copy the current stack pointer to RBP  
  
    # function body  
    ...  
  
    #epilogue  
    movq %rbp, %rsp # retrieve the original RSP value  
    popq %rbp       # restore the original RBP value  
    ret
```

- The prologue saves the caller's value for `%rbp` into the callee's stack frame and the current stack pointer in `%rbp`
- This adjusted value of `%rbp` is the callee's frame pointer
 - The callee will not change this value until just before it returns
 - The frame pointer provides a stable reference point for local variables and caller arguments
- The epilogue ensures that the caller's `%rsp` and `%rbp` are restored, which deallocates all callee's reserved temporary space



- Before the function (with more than 6 arguments) is called, the caller places the 7+ required input values in the stack
- When the called function returns, those values are still on the stack (since the callee function accessed them without popping them off of the stack)
- The caller should remove the old input values from the stack to get the stack back to where it was before the function call
 - This is particularly important if other values are going to be popped out of the stack after the callee's return
 - Ensures that the `ret` instruction at the end of the caller finds the proper return value on top of the stack
- While you can use the `POP` instruction to do this, you can also just move the RSP stack pointer back to the original location before the function call

- Adding back the size of the data elements pushed onto the stack using the ADD instruction clears all the pushed data with a single instruction

Cleaning the Stack after Return

```
function:  
...  
  
pushq %rax      # push 8th parameter onto stack  
pushq %rbx      # push 7th parameter onto stack  
call utilfunc  
addq $16, %rsp # get the stack back to where it was before the function call  
  
...  
ret
```

- Consider the following functions:

Function f1

```
long f1(long a, long b, long c,
         long d, long e, long f){
    return a + f2(a*b,b,c,d,e,f,10,-20);
}
```

Function f2

```
long f2(long a, long b, long c,
         long d, long e, long f,
         long g, long h){
    /* do something with parameters */
    ...
    return h/g;
}
```

- How can you implement an equivalent code in Assembly?

Passing Parameters

```
#long f1(a,b,c,d,e,f)
f1:
# a in %rdi, b in %rsi, c in %rdx
# d in %rcx, e in %r8, f in %r9

pushq %rdi          # save %rdi on stack

imulq %rsi, %rdi    # a * b
# prepare function call
pushq $-20           # 8th parameter
pushq $10             # 7th parameter
call f2               # f2(a*b,b,c,d,e,f,10,-20);
addq $16, %rsp        # clear 7th and 8th parameters

popq %rdi            # restore %rdi
addq %rdi, %rax      # a + f2(...)

ret
```

Accessing Parameters

```
#long f2(a,b,c,d,e,f,g,h)
f2:
    pushq %rbp          #prologue
    movq %rsp, %rbp

    # a in %rdi, b in %rsi, c in %rdx
    # d in %rcx, e in %r8, f in %r9
    # g in 16(%rbp), h in 24(%rbp)

    pushq %rbx           # save %rbx

    # do something with parameters
    ...
    movq 16(%rbp), %rbx      # get g from stack
    movq 24(%rbp), %rax      # get h from stack
    cqto
    idivq %rbx             # h/g

    popq %rbx              # restore %rbx

    movq %rbp, %rsp          # epilogue
    popq %rbp
    ret
```



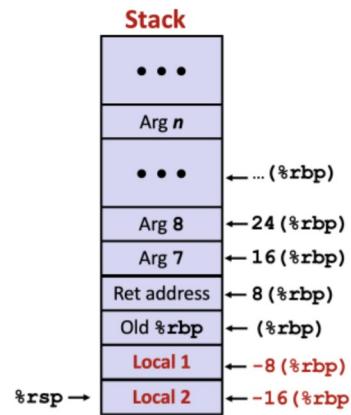
- At times, local data must be stored in memory. Common cases of this include:
 - There are not enough registers to hold all of the local data
 - The address operator '&' is applied to a local variable, and hence we must be able to generate an address for it
 - Some of the local variables are arrays or structures and hence must be accessed by array or structure references
- When looking for an easy location for working storage for data elements within a function, the stack again comes to the rescue

- The idea is to allocate storage for local variables when entering the function and release it before returning
- This space should be reserved after setting the RBP register to point to the top of the stack without affecting how the function's parameters are accessed
- Typically, a function allocates space on the stack frame by decrementing the stack pointer (%rsp)

Reserving Space for Local Variable

```
function:  
# prologue  
pushq %rbp  
movq %rsp, %rbp  
subq $16, %rsp    # reserve 16 bytes for local variables  
# this space can be used, for e.g., as either  
# four 32-bit values or two 64-bit values  
...
```

- By keeping %rbp pointing to the initial position throughout the execution of the function, we can refer local variables and parameters with fixed offsets relative to %rbp



- Now, if any data is pushed onto the stack, it will be placed after the local variables
- Local variables are then preserved and they can still be accessed via RBP
- When the epilogue of the function is reached and the RSP register is set back to its original value, the local variables will be lost from the stack
- They will not be directly accessible using the RSP or RBP registers from the calling function (thus the term "local variables")

- Consider the following functions:

Function swap_add

```
long swap_add(long *xp, long *yp){  
    long x = *xp;  
    long y = *yp;  
    *xp = y;  
    *yp = x;  
    return x + y;  
}
```

Function f1

```
long f1(){  
    long arg1 = 100;  
    long arg2 = 200;  
    long sum = swap_add(&arg1,&arg2);  
    long diff = arg1 - arg2;  
    return sum * diff;  
}
```

- How can you implement an equivalent code in Assembly?

Function swap_add

```
#long swap_add(long *xp, long *yp)
swap_add:
    movq (%rdi), %rax      # get x
    movq (%rsi), %rcx      # get y
    movq %rcx, (%rdi)      # *xp = y
    movq %rax, (%rsi)      # *yp = x
    addq %rcx, %rax        # x + y
    ret
```

Function f1

```
#long f1()
f1:
    pushq %rbp          # prologue
    movq %rsp, %rbp
    subq $16, %rsp        # allocate 16 bytes for local variables

    movq $100, -8(%rbp)   # arg1 = 100;
    movq $200, -16(%rbp)  # arg2 = 200;

    leaq -8(%rbp), %rdi   # prepare function call
    leaq -16(%rbp), %rsi  # compute &arg1 as first argument
    call swap_add          # compute &arg2 as second argument
    # call swap_add(&arg1,&arg2)

    movq -8(%rbp), %rdx   # get arg1
    subq -16(%rbp), %rdx  # diff = arg1 - arg2
    imulq %rdx             # sum * diff

    movq %rbp, %rsp        # epilogue (deallocates 16 bytes)
    popq %rbp
    ret
```

