

Computer Architecture (Practical Class)

C and Assembly: Bit-level Operations

Luís Nogueira

Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto

lmn@isep.ipp.pt

2023/2024

Bit-level operations modify one or more bits at a time and are used to manipulate binary numbers (stored in memory or registers)

- Are very efficient operations, directly supported by the processor
- Can be used to:
 - Extract information from groups of bits
 - Perform multiplications and divisions by powers of two several times faster
 - Apply an operation to multiple data within a single variable (if the variables has groups of bits representing different data)

Two major groups of bit-level operations

- Boolean logic
 - Each bit is compared individually using the logic function specified
 - Logic functions: AND, OR, XOR, NOT
- Bit movements
 - Shift bits left/right (multiply/divide by powers of two)
 - Rotate bits left or right

AND (the “ & ” operator in C)

X	Y	AND
0	0	0
0	1	0
1	0	0
1	1	1

- Usage: AND *origin, destination*

- operation: *destination* = *destination* AND *origin* (the result is placed in *destination*)
- origin* can be a memory address, a constant value or a register
- destination* can be a memory address or a register
- the AND instruction can operate on numbers of 8(b), 16(w), 32(l), or 64(q) bits

OR (the “ | ” operator in C)

X	Y	OR
0	0	0
0	1	1
1	0	1
1	1	1

- Usage: `OR origin, destination`
 - operation: `destination = destination OR origin` (the result is placed in `destination`)
 - `origin` can be a memory address, a constant value or a register
 - `destination` can be a memory address or a register
 - the OR instruction can operate on numbers of 8(b), 16(w), 32(l), or 64(q) bits

XOR (the “`^`” operator in C)

X	Y	XOR
0	0	0
0	1	1
1	0	1
1	1	0

- Usage: `XOR origin, destination`
 - operation: `destination = destination XOR origin` (the result is placed in `destination`)
 - `origin` can be a memory address, a constant value or a register
 - `destination` can be a memory address or a register
 - the XOR instruction can operate on numbers of 8(b), 16(w), 32(l), or 64(q) bits

NOT (the “~” operator in C)

X	NOT
0	1
1	0

• Usage: NOT *destination*

- operation: *destination* = NOT *destination* (the result is placed in *destination*)
- destination* can be a memory address or a register
- the NOT instruction can operate on numbers of 8(b), 16(w), 32(l), or 64(q) bits

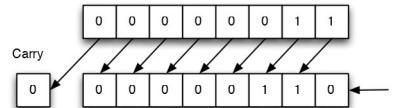
Important notes

- The Assembly instruction NEG changes the sign of an integer
- The logic C operator “!” is not a bit-level operator; it considers the entire number as a logical value

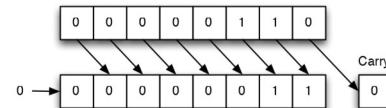
y	NOT y	NEG y	!y
0	-1	0	1
-1	0	1	0
1	-2	-1	0

- Shifting bits left and right is a very easy operation to implement on a processor. It is often used as a fast way to implement multiplication/division
- Consider a binary number:

- Shifting a digit left (entering a zero) corresponds to a multiplication by 2



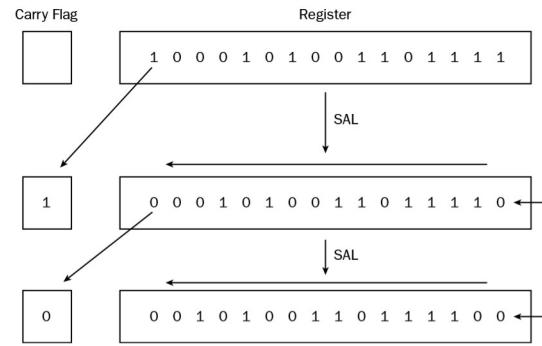
- Shifting a digit right (losing the rightmost digit) corresponds to a division by 2



- This is true for any base b (shifting a digit left/right corresponds to multiply/divide by b)

SHL/SAL (the “`<<`” operator in C)

- SHL and SAL are equivalent operations
 - They exist for consistency with the right shift (we will see why in a moment)
- Shifts bits to the left; zeros enter on the right and the last bit to exit left goes to the carry flag

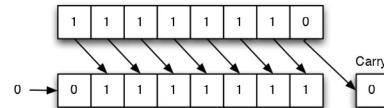


SHL/SAL - Three formats

- SHL *destination* (or SAL *destination*)
 - Shifts the *destination* value left one position (equivalent to *destination* * = 2)
- SHL %cl, *destination* (or SAL %cl, *destination*)
 - Shifts the *destination* value left by the number of times specified in the CL register (equivalent to *destination* * = 2^{CL})
- SHL \$n, *destination* (or SAL \$n, *destination*)
 - Shifts the *destination* value left by the number of times specified by a constant value *n* (equivalent to *destination* * = 2^n)
- In all formats, *destination* can be a memory address or a register
- The SHL/SAL instructions can operate on numbers of 8(b), 16(w), 32(l), or 64(q) bits

Notes about signed numbers

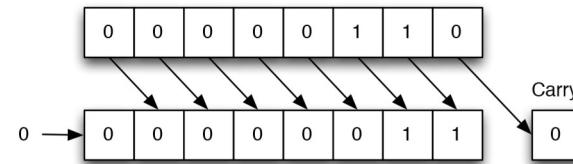
- Performing a right shift on a signed integer value may adversely affect the sign of the integer
- When shifted to the right, a negative number will lose its sign if we zero-fill the leading bits



- To solve this problem there is a distinction between the right-shift instructions:
 - SHR - the **logic shift** to the right *does not preserve the signal*;
 - SAR - the **arithmetic shift** to the right *preserves the signal*.

SHR (the “`>>`” operator in C, when applied to unsigned integers)

- Logic shift to the right
- Shifts bits to the right; zeros enter on the left and the last bit to exit to the right goes to the carry flag (similar to the left shift)
- Therefore, *does not preserve the signal* of the number

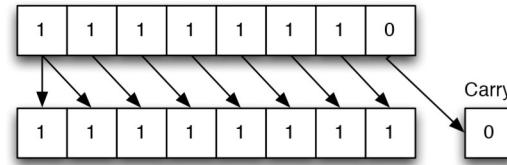


SHR - Three formats

- SHR *destination*
 - Shifts the *destination* value right one position (equivalent to *destination* / = 2).
- SHR %*cl*, *destination*
 - Shifts the *destination* value right by the number of times specified in the CL register (equivalent to *destination* / = 2^{CL}).
- SHR \$*n*, *destination*
 - Shifts the *destination* value right by the number of times specified by a constant value *n* (equivalent to *destination* / = 2^n)
- In all formats, *destination* can be a memory address or a register
- The SHR instruction can operate on numbers of 8(b), 16(w), 32(l), or 64(q) bits

SAR (the “`>>`” operator in C, when applied to signed integers)

- Arithmetic shift to the right
- Shifts bits to the right; either clears or sets the bits entered on the left, according to the sign of the integer. The last bit that exits to the right goes to the carry flag
- Therefore, *preserves the signal* of the number

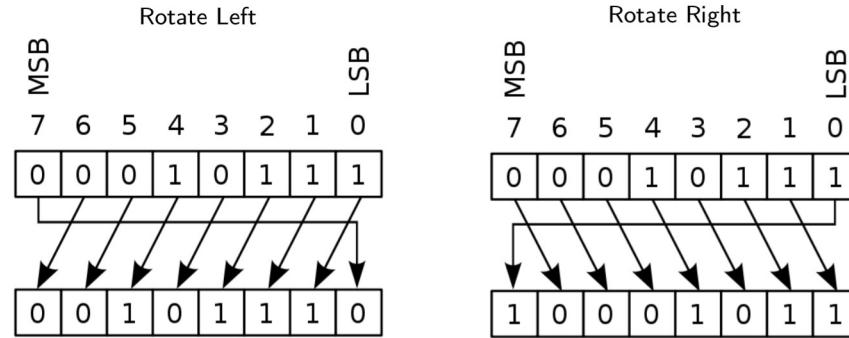


SAR - Three formats (similar to SHR):

- SAR *destination*
- SAR %cl, *destination*
- SAR \$n, *destination*
- In all formats, *destination* can be a memory address or a register
- The SAR instruction can operate on numbers of 8(b), 16(w), 32(l), or 64(q) bits

ROL/ROR (no equivalent operation in C)

- ROL - Bit rotation to the left
 - ROR - Bit rotation to the right
 - Perform just like the shift instructions, except the overflow bits are pushed back into the other end of the value instead of being dropped.

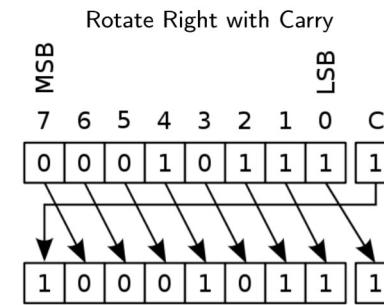
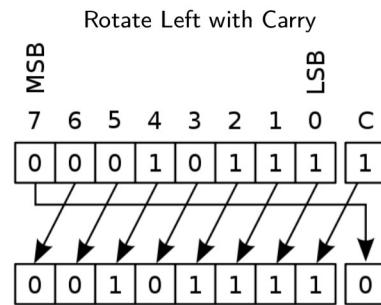


ROL/ROR - Three formats (similar to shift instructions):

- $RO\{L/R\} \ destination$
- $RO\{L/R\} \%cl, destination$
- $RO\{L/R\} \$n, destination$
- In all formats, *destination* can be a memory address or a register
- The ROL/ROR instructions can operate on numbers of 8(b), 16(w), 32(l), or 64(q) bits

RCL/RCR (no equivalent operation in C)

- RCL - Bit rotation to the left *with carry*
- RCR - Bit rotation to the right *with carry*
- Perform bit rotation, but the first entering bit comes from carry and overflow bits go to carry.



RCL/RCR - Three formats (similar to ROL/ROR):

- $RC\{L/R\} \ destination$
- $RC\{L/R\} \%cl, destination$
- $RC\{L/R\} \$n, destination$
- In all formats, *destination* can be a memory address or a register
- The RCL/RCR instructions can operate on numbers of 8(b), 16(w), 32(l), or 64(q) bits

Bit masks

- One common use of bit-level operations is to implement *masking* operations
- A *mask* is a bit pattern that indicates a selected set of bits within a number

For example, assume that %al has a binary value of 00101100

- How can we get the value of the 4 least significant bits?

To solve the problem we need to:

- ① Determine the logic operator to use
- ② Determine the mask, based on the selected operator

```
movb $0b00101100, %al  
movb $0b00001111, %ah  
andb %ah, %al      # %al = 0b00001100
```

Assume now we want to replace the 4 least significant bits of %al for the 4 least significant bits of %cl. We need to perform three steps:

- ① Set the 4 least significant bits of %al to zero using the and operation
- ② Set the 4 most significant bits of %cl to zero using the and operation
- ③ Replace the bits in %al with the bits in %cl using the or operation

```
movb $0b00101100, %al
movb $0b01000011, %cl

movb $0b11110000, %ah
andb %ah, %al          # %al = 0b00100000
notb %ah               # %ah = 0b00001111 (inverts the mask)
andb %ah, %cl          # %cl = 0b00000011

orb %cl, %al           # %al = 0b00100011
```

- The `xor` operation can be used to set a number to zero or compare two numbers
 - The result of a `xor` operation between two equal numbers is zero

```
xorl %ebx, %ecx  
jz is_equal
```

- The `xor` operation is also used to invert the bits of a register
 - Placing 0/1 on the mask to keep/invert the original bit

```
movb $0b00101100, %al  
movb $0b00001111, %ah  
xorb %ah, %al      # %al = 0b00100011
```

- The **AND** operator can be used to:
 - Set bits to zero - use a bit mask with those bits set to zero
 - Get the value of a few bits - use a bit mask with those bits set to one
 - "Round" a number to a power of 2 (2^x) - use a bit mask with x least significant bits to zero
 - Get the remainder of the division by a power of two (2^x) - use a bit mask with x least significant bits to one
- The **OR** operator can be used to:
 - Set bits to one - use a bit mask with those bits set to one
 - Join the bits of two numbers - perform an **or** between the two numbers
- The **XOR** operator can be used to:
 - Check if two numbers are equal - if a **xor** between the two numbers is zero
 - Invert bits - use a bit mask with those bits set to one

```
#define mask(n)      ((1<<(n))-1)
#define mask2(n1,n2)  (mask(n2-n1+1)<<(n1))

#include <stdio.h>
int main ()
{
    printf("%d\n", mask(2));
    printf("%d\n", mask(3));
    printf("%d\n", mask2(1,2));
    printf("%d\n", mask2(1,3));
    return 0;
}
```

- The above code prints...

- A. the numbers 2, 3, 2, 3
- B. the numbers 4, 8, 6, 9
- C. the numbers 3, 7, 6, 14
- D. None of the above.

Implement in C and in Assembly the function:

```
void sum_bytes(int a, int b, char *sum)
```

- Sums the most significant byte of the integer a with the least significant byte of b
- Writes the result in the memory address pointer by sum

Function sum_bytes in C

```
void sum_bytes(int a, int b, char *sum){  
    char msb_a = (a & 0xFF000000) >> 24;  
    char lsb_b = b & 0x000000FF;  
  
    *sum = msb_a + lsb_b;  
}
```

Function sum_bytes in Assembly

```
sum_bytes:  
# a in %edi, b in %esi, *sum in %rdx  
  
    shr $24, %edi      # MSB in %dil  
    addb %dil, %sil  
    movb %sil, (%rdx)  # *sum = %dil + %sil  
    ret
```

Given the following C code to multiply two numbers using logic and bit movement operations, implement the equivalent in Assembly

Function multiply in C

```
int multiply(int a, int b){  
    int res = 0;  
  
    while (b != 0){  
        if ((b & 1) == 1)  
            res = res + a;  
        a = a << 1;  
        b = (unsigned int)b >> 1;  
    }  
  
    return res;  
}
```

Function multiply in Assembly

```
#int mult(int a, int b)
mult:
    # a in %edi, b in %esi
    movl $0, %eax      # res = 0

loop_mult:
    cmpl $0, %esi      # while(b!=0)
    je end
    pushq %rsi          # save b in stack
    andl $1, %esi       # b = b & 1
    cmpl $1, %esi
    jne next
    addl %edi, %eax    # res = res + a
next:
    popq %rsi          # get b from stack
    shll %edi          # a << 1
    shr1 %esi          # b >> 1
    jmp loop_mult
end:
    ret
```