

# Integer arithmetic

Arquitectura de Computadores  
Departamento de Engenharia Informática  
Instituto Superior de Engenharia do Porto

Luís Nogueira ([lmn@isep.ipp.pt](mailto:lmn@isep.ipp.pt))

# Today: Integer arithmetic

- Encoding integers
- Addition
- Negation
- Multiplication
- Shifting

# Integer arithmetic

- Many beginning programmers are surprised to find that...
  - adding two positive numbers can yield a negative result
  - the comparison  $x < y$  can yield a different result than the comparison  $x - y < 0$
- These properties are artifacts of the finite nature of computer arithmetic
  - As we will see, the integer arithmetic performed by computers is really a form of modular arithmetic
- Understanding the nuances of computer arithmetic can help programmers write more reliable code

# Today: Integer arithmetic

- Addition
- Negation
- Multiplication
- Shifting

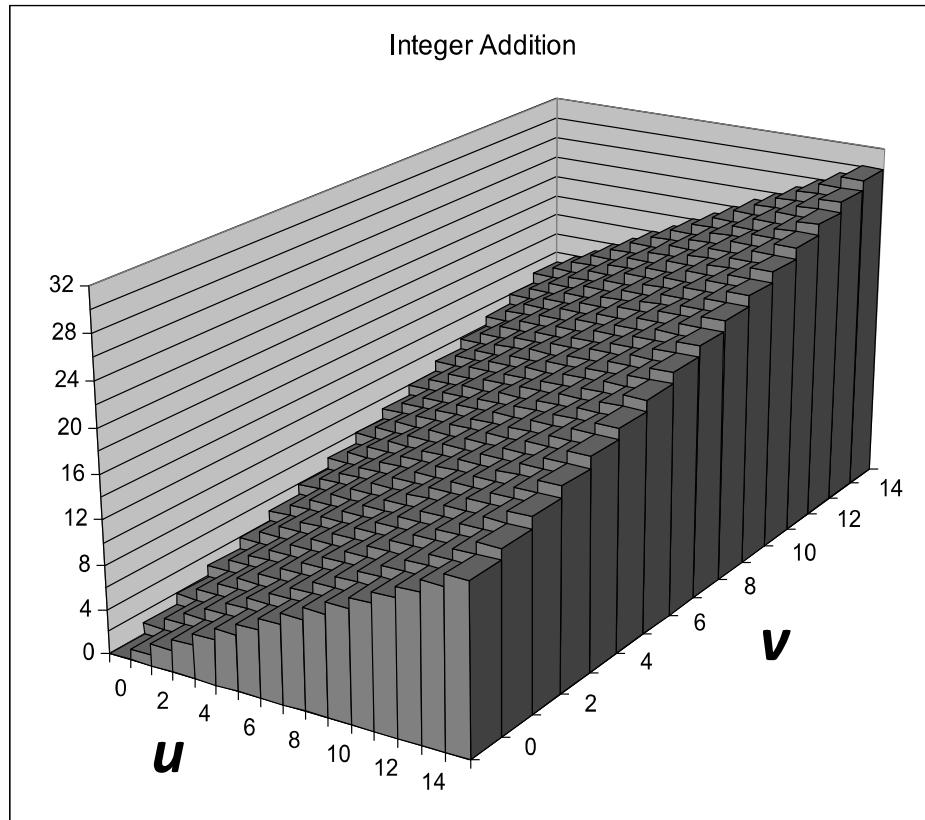
# Visualizing (mathematical) integer addition

## ■ Example:

- 4-bit integers  $u, v$
- Compute true sum with  $\text{Add}_4(u, v)$

## ■ Values increase linearly with $u$ and $v$

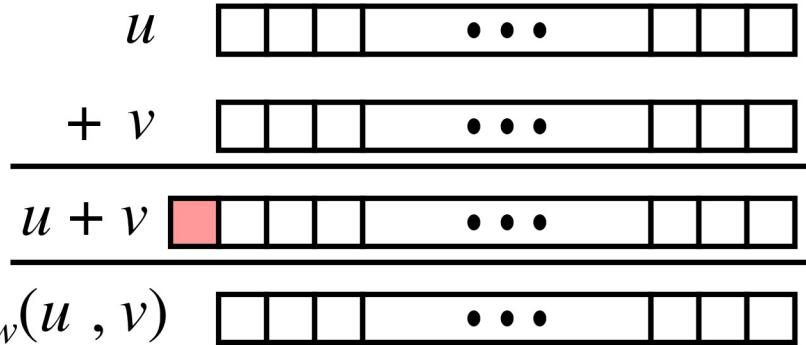
- Forms planar surface



$$\text{Add}_4(u, v) = u + v$$

# Unsigned addition in C

Operands:  $w$  bits



True sum:  $w+1$  bits

Discard carry:  $w$  bits

- Unsigned addition ignores carry output

- Implements modular arithmetic

$$s = \text{UAdd}_w(u, v) = (u + v) \bmod 2^w$$

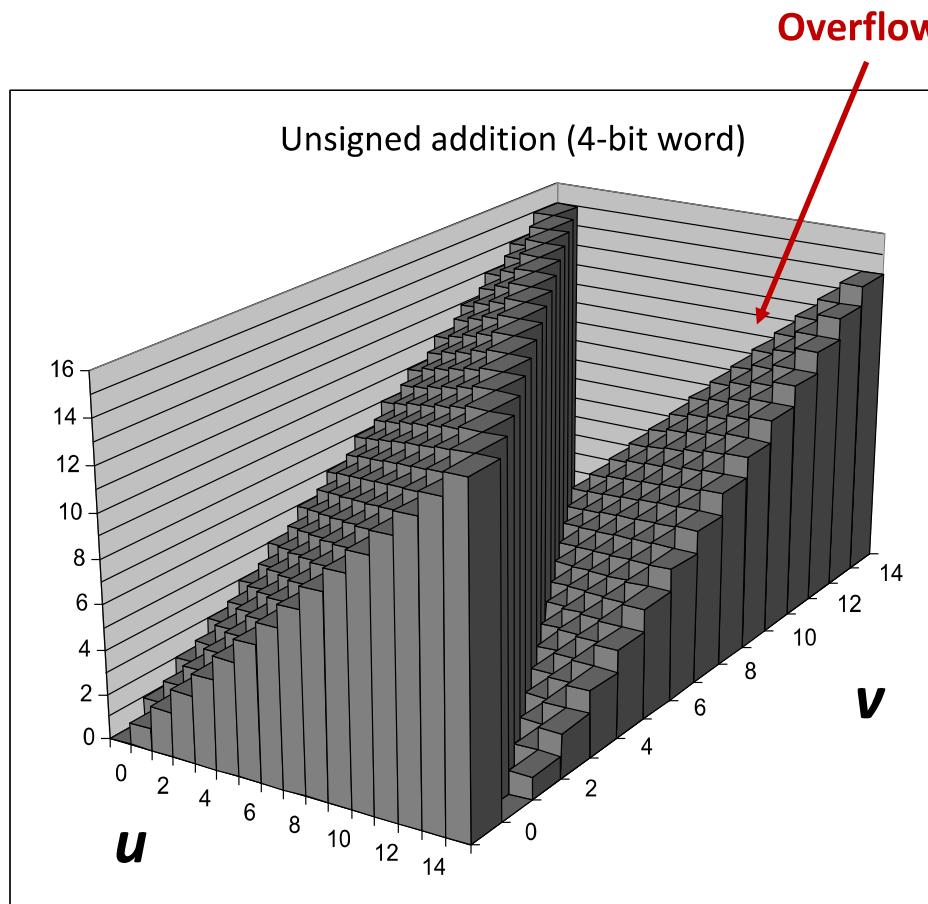
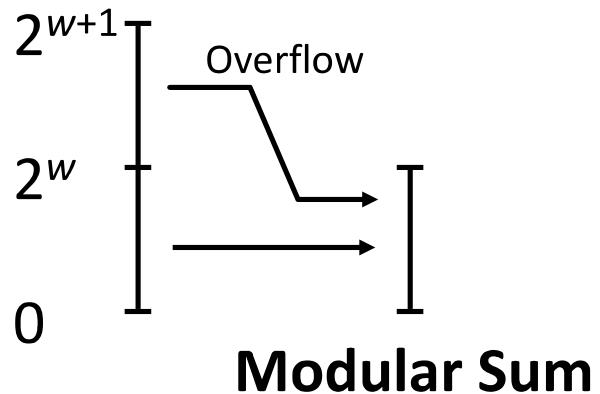
$$UAdd_w(u, v) = \begin{cases} u + v, & u + v < 2^w \\ u + v - 2^w, & u + v \geq 2^w \end{cases}$$

# Visualizing unsigned addition

## Wraps around

- If true sum  $\geq 2^w$
- At most once

**True Sum**



$$\text{UAdd}_4(u, v) = (u + v) \bmod 16$$

# Unsigned addition

<b>unsigned char</b>	1110 1001	E9	223
+	1101 0101	+ D5	+ 213
	<b>1 1011 1110</b>	<b>1BE</b>	<b>446</b>
	<b>1011 1110</b>	<b>BE</b>	<b>190</b>

Hex    Decimal  
Binary

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

# Two's complement addition in C

Operands:  $w$  bits

$$\begin{array}{r}
 u \quad \boxed{\phantom{0}} \boxed{\phantom{0}} \boxed{\phantom{0}} \dots \boxed{\phantom{0}} \boxed{\phantom{0}} \\
 + \quad v \quad \boxed{\phantom{0}} \boxed{\phantom{0}} \boxed{\phantom{0}} \dots \boxed{\phantom{0}} \boxed{\phantom{0}} \\
 \hline
 u + v \quad \boxed{\textcolor{red}{1}} \boxed{\phantom{0}} \boxed{\phantom{0}} \boxed{\phantom{0}} \dots \boxed{\phantom{0}} \boxed{\phantom{0}}
 \end{array}$$

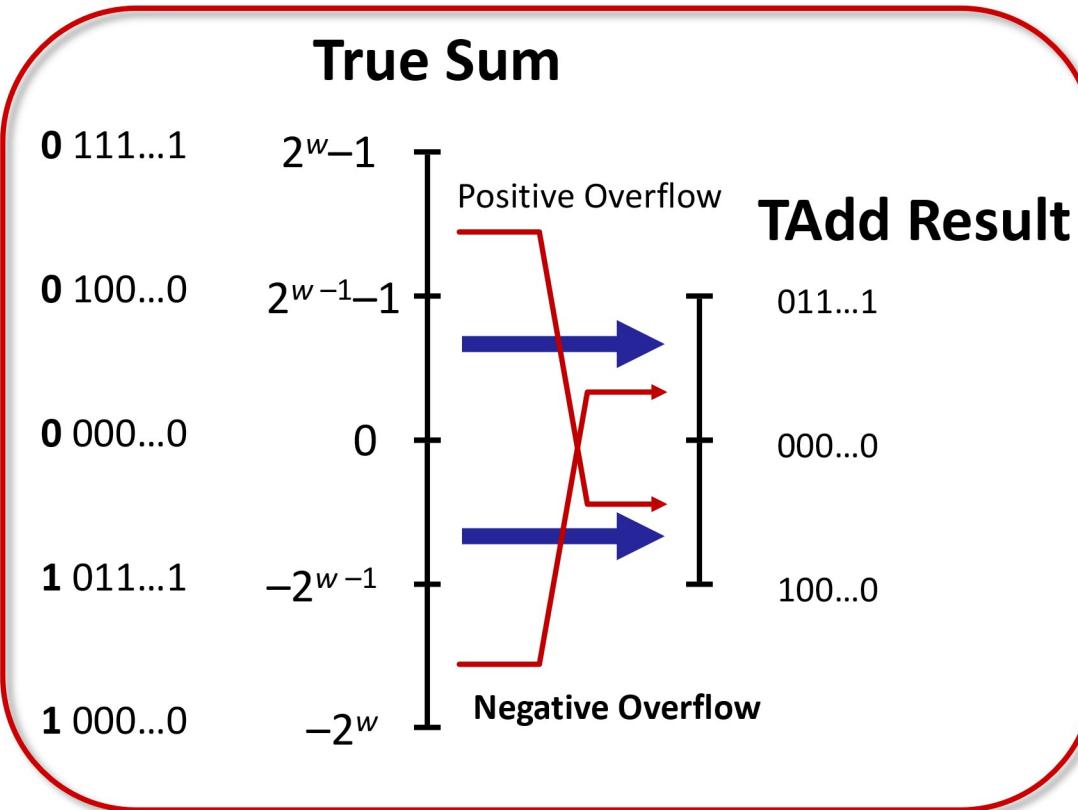
True sum:  $w+1$  bits

Discard carry:  $w$  bits       $\text{TAdd}_w(u, v)$        $\boxed{\phantom{0}} \boxed{\phantom{0}} \boxed{\phantom{0}} \dots \boxed{\phantom{0}} \boxed{\phantom{0}}$

- As before, we avoid ever-expanding data sizes by truncating the representation to  $w$  bits
- However, we must decide what to do when the result is either too large (positive) or too small (negative) to represent

# Two's complement addition overflow

- Drop off most significant bit
- Treat remaining bits as two's complement integer



$$TAdd_w = \begin{cases} u + v + 2^{w-1}, & u + v < TMin_w \\ u + v, & TMin_w \leq u + v \leq TMax_w \\ u + v - 2^{w-1}, & TMax_w < u + v \end{cases}$$

(Negative Overflow)

(Positive Overflow)

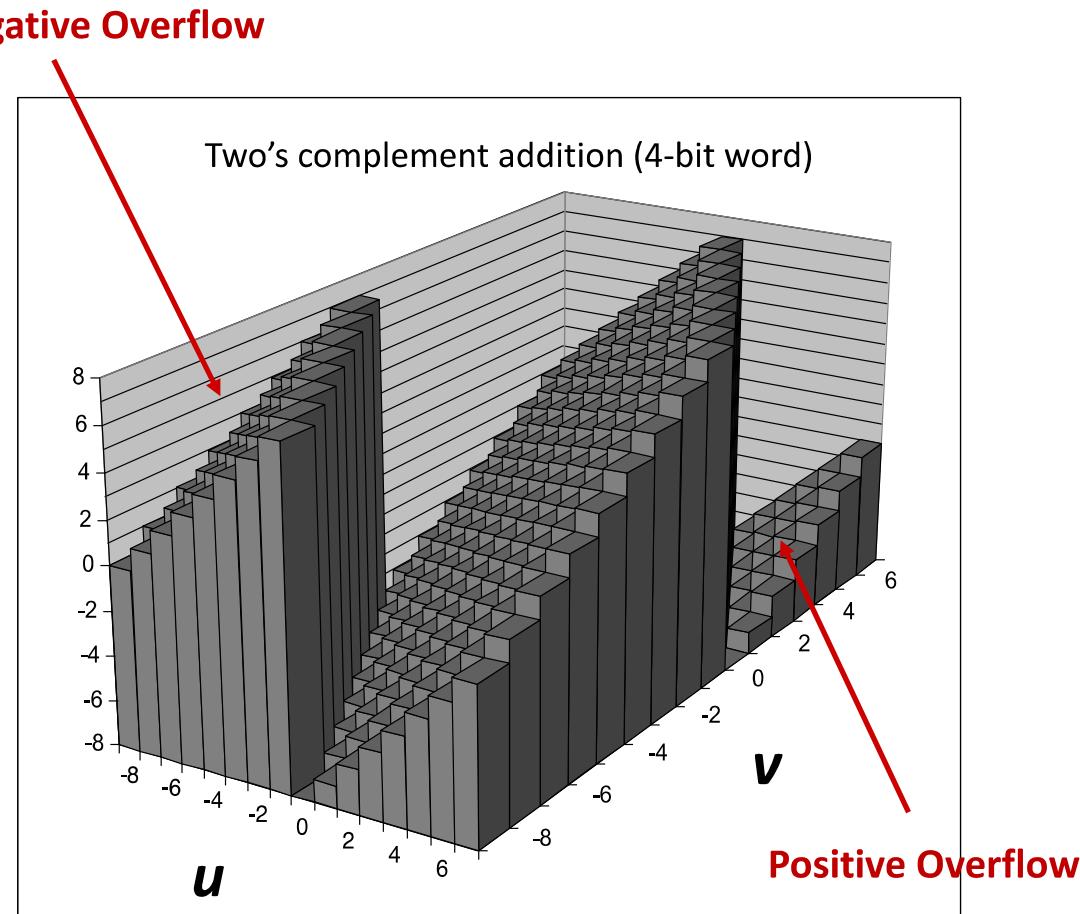
# Visualizing two's complement addition

## ■ Example:

- 4-bit two's complement
- Range from -8 to +7

## ■ Wraps Around

- If  $\text{sum} \geq 2^{w-1}$ 
  - Becomes negative
  - At most once
- If  $\text{sum} < -2^{w-1}$ 
  - Becomes positive
  - At most once



$$\text{TAdd}_4(u, v) = \text{U2T}_4[(x+y) \bmod 16]$$

# Two's complement addition

**char**

$$\begin{array}{r}
 & 1110 \ 1001 & E9 & 223 \\
 + & 1101 \ 0101 & + D5 & + 213 \\
 \hline
 1 & 1011 \ 1110 & 1BE & 446 \\
 \hline
 & 1011 \ 1110 & BE & 190
 \end{array}$$

Hex    Decimal  
Binary

Hex	Decimal	Binary
0	0	0000
1	1	0001
2	2	0010
3	3	0011
4	4	0100
5	5	0101
6	6	0110
7	7	0111
8	8	1000
9	9	1001
A	10	1010
B	11	1011
C	12	1100
D	13	1101
E	14	1110
F	15	1111

- Signed and unsigned addition have identical bit-level behavior

# Practice problem

- Write a function with the following prototype:

```
int add_ok(int x, int y);
```

- This function should return 1 if arguments x and y can be added without causing overflow

# Practice problem

```
int add_ok(int x, int y) {  
    int sum = x+y;  
    int neg_over = x < 0 && y < 0 && sum >= 0;  
    int pos_over = x >= 0 && y >= 0 && sum < 0;  
  
    return !neg_over && !pos_over;  
}
```

- In C, overflows are not signaled as errors
- We can check if an overflow has occurred on  $x + y$  by seeing, if and only if,:
  - $sum < x$  (or equivalently,  $sum < y$ ) for unsigned addition
  - The above conditions for signed addition

# Summary: Signed vs Unsigned addition

- TAdd and UAdd have identical bit-level behavior
  - Most CPUs use the same machine instruction to perform either unsigned or signed addition
- Signed vs. Unsigned addition in C

```
int s, t, u, v;  
s = (int) ((unsigned) u + (unsigned) v);  
t = u + v
```

Will give  $s == t$

# Today: Integer arithmetic

- Addition
- Negation
- Multiplication
- Shifting

# Two's complement negation

- Find the corresponding negative/positive number with the same absolute value
- Every number  $x$  in the range  $-2^{w-1} \leq x < 2^{w-1}$  has an additive inverse
  - For  $x \neq -2^{w-1}$ , we can see that its additive inverse is simply  $-x$
- For  $x = -2^{w-1} = \text{TMin}_w$ ,  $-x = 2^w$  which cannot be represented as a  $w$ -bit number

$$x = \begin{cases} -2^{w-1}, & x = -2^{w-1} \\ -x, & x > -2^{w-1} \end{cases}$$

# Two's complement negation

## ■ Solution: complement and increment

$$\sim x + 1 == -x$$

## ■ Observation:

$$\sim x + x == 1111\dots111 == -1$$

$$\begin{array}{r} x \quad \boxed{1} \boxed{0} \boxed{0} \boxed{1} \boxed{1} \boxed{1} \boxed{0} \boxed{1} \\ + \quad \sim x \quad \boxed{0} \boxed{1} \boxed{1} \boxed{0} \boxed{0} \boxed{0} \boxed{1} \boxed{0} \\ \hline -1 \quad \boxed{1} \boxed{1} \boxed{1} \boxed{1} \boxed{1} \boxed{1} \boxed{1} \boxed{1} \end{array}$$

# Examples: Complement & Increment

**x = 15213**

	<b>Decimal</b>	<b>Hex</b>	<b>Binary</b>
<b>x</b>	15 213	3B 6D	00111011 01101101
<b>~x</b>	-15 214	C4 92	11000100 10010010
<b>~x+1</b>	-15 213	C4 93	11000100 10010011

**x = 0**

	<b>Decimal</b>	<b>Hex</b>	<b>Binary</b>
<b>0</b>	0	00 00	00000000 00000000
<b>~0</b>	-1	FF FF	11111111 11111111
<b>~0+1</b>	0	00 00	00000000 00000000

# Today: Integer arithmetic

- Addition
- Negation
- Multiplication
- Shifting

# Multiplication

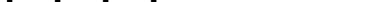
- Computing product of  $w$ -bit numbers  $x$  and  $y$ 
  - Either signed or unsigned
- Exact results can require as many as  $2^*w$  bits to represent
  - Most cases would fit into  $2^*w - 1$  bits, but the special case of  $2^{2w-2}$  requires the full  $2^*w$  bits (to include a sign bit of 0)
- So, maintaining exact results...
  - Would need to keep expanding word size with each product computed
  - Is done in software, if needed (e.g., by “arbitrary precision” arithmetic packages)

# Unsigned multiplication in C

## Operands: $w$ bits

$$* \quad v \quad \boxed{\phantom{0} \phantom{0} \phantom{0}} \quad \dots \quad \boxed{\phantom{0} \phantom{0}}$$

**True product:  $2^*w$  bits**  $u \cdot v$  

**Truncate:**  $w$  bits       $\text{UMult}_w(u, v)$  

## ■ Standard multiplication algorithm

- Truncate result to w-bit number

## ■ Implements modular arithmetic

$$\text{UMult}_w(u, v) = (u \cdot v) \bmod 2^w$$

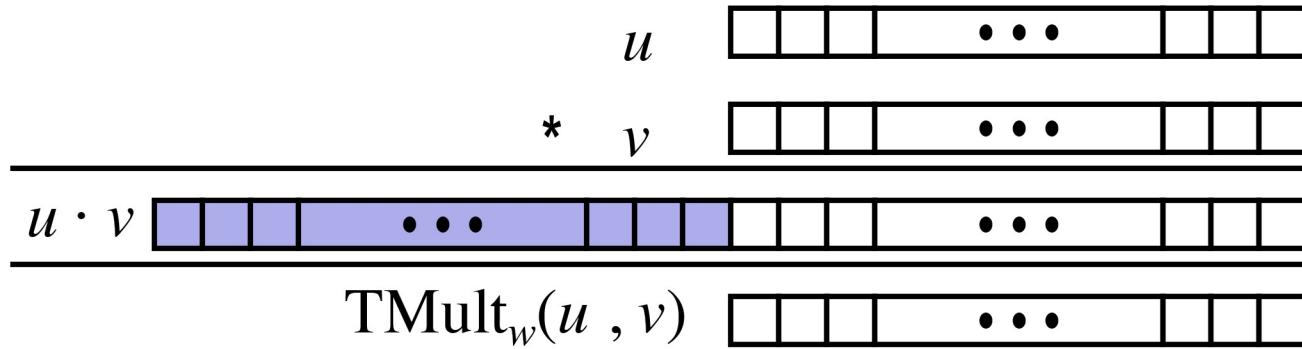
# Unsigned multiplication in C

unsigned char

$$\begin{array}{r} & \begin{array}{r} 1110 \ 1001 \\ * \quad \quad \quad 1101 \ 0101 \\ \hline 1100 \ 0001 \ 1101 \ 1101 \\ \hline 1101 \ 1101 \end{array} & \begin{array}{r} E9 \\ * \quad \quad \quad D5 \\ \hline C1DD \\ DD \end{array} & \begin{array}{r} 233 \\ * \quad \quad \quad 213 \\ \hline 49629 \\ 221 \end{array} \end{array}$$

# Two's complement multiplication in C

Operands:  $w$  bits



- Compute exact product and ignores high order bits
  - Some of which are different for signed vs. unsigned multiplication
  
- Lower bits are the same of unsigned multiplication
  - Treat them as a two's complement integer

# Signed multiplication in C

char

$$\begin{array}{r} & 1110 \ 1001 & E9 & -23 \\ * & 1101 \ 0101 & * D5 & * -43 \\ \hline 0000 \ 0011 \ 1101 \ 1101 & \hline 03DD & \hline 989 \\ \hline 1101 \ 1101 & DD & -35 \end{array}$$

# Signed vs Unsigned multiplication

- TMul and UMul have identical bit-level behavior for the low-order  $w$  bits, even though the full  $2w$ -bit differ
  - Separate instructions are provided in x86-64 for signed and unsigned multiplication
- Multiplication in C is performed by truncating the  $2w$ -bit product to  $w$  bits

```
int x, y;  
unsigned ux = (unsigned) x;  
unsigned uy = (unsigned) y;  
  
unsigned up = ux * uy;  
int p = x * y;
```

up == (unsigned) p

# Today: Integer arithmetic

- Addition
- Negation
- Multiplication
- Shifting

# Multiplying by constants

## ■ Most machines shift and add faster than multiply

- The integer multiply instruction is fairly slow, requiring 10 or more clock cycles
- Other integer operations—such as addition, subtraction, bit-level operations, and shifting—require only 1 clock cycle

## ■ Compiler optimization

- Replace multiplications by constant factors with combinations of shift and addition/subtraction operations

## ■ Examples

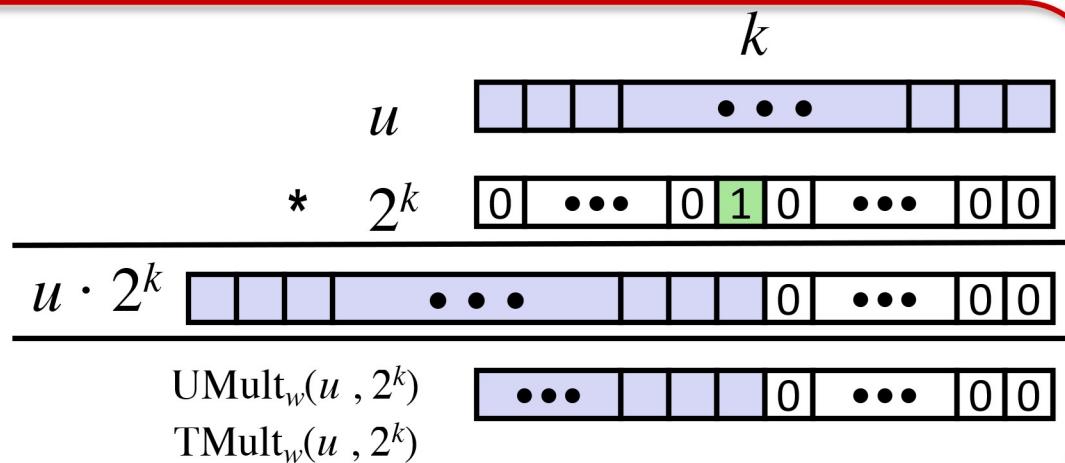
- `u << 3 == u * 8`
- `(u << 5) - (u << 3) == u * 24`

# Power-of-2 multiply with shift

**Operands:**  $w$  bits

**True product:**  $w+k$  bits

**Truncate:**  $w$  bits



## ■ Operation

- $u \ll k$  gives  $u * 2^k$

## ■ Yield the same result

- Both signed and unsigned
- Even in overflow

# Compiled multiplication code

## C function

```
long mul_12(long x) {  
    return x*12;  
}
```

## Compiled arithmetic operations

```
movq %rdi, %rax  
shlq $3, %rax  
shlq $2, %rdi  
addq %rdi, %rax
```

## Explanation

```
t1 <- x * 8;  
t2 <- x * 4;  
return t1 + t2;
```

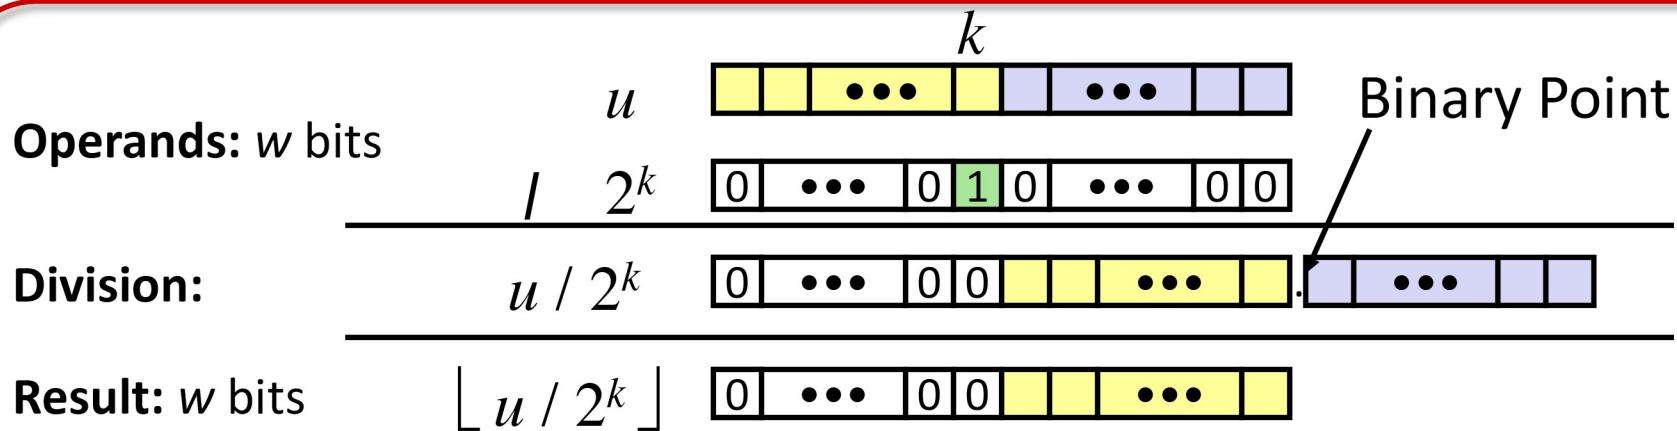
- C compiler automatically generates shift/add code when multiplying by constant

- Always uses logical shift

# Power-of-2 divide with shift

- Integer division on most machines is even slower than integer multiplication
  - Requiring 30 or more clock cycles
- Dividing by a power of 2 can also be performed using shift operations
  - Uses a right shift rather than a left shift
- Unsigned vs Two's complement
  - Unsigned: logical shift
  - Two's complement: arithmetic shift

# Unsigned power-of-2 divide with shift



- Quotient of unsigned by power-of-2

- $u \gg k$  gives  $\lfloor u / 2^k \rfloor$

- Uses logical shift

# Unsigned power-of-2 divide with shift

	Division	Computed	Hex	Binary
x	15 213	15 213	3B 6D	00111011 01101101
x >> 1	7 606.5	7 606	1D B6	00011101 10110110
x >> 4	950.8125	950	03 B6	00000011 10110110
x >> 8	59.4257813	59	00 3B	00000000 00111011

- The result of shifting consistently rounds toward zero
  - As is the convention for integer division

# Compiled unsigned division code

## C function

```
unsigned long udiv_8(unsigned long x) {  
    return x/8;  
}
```

## Compiled arithmetic operations

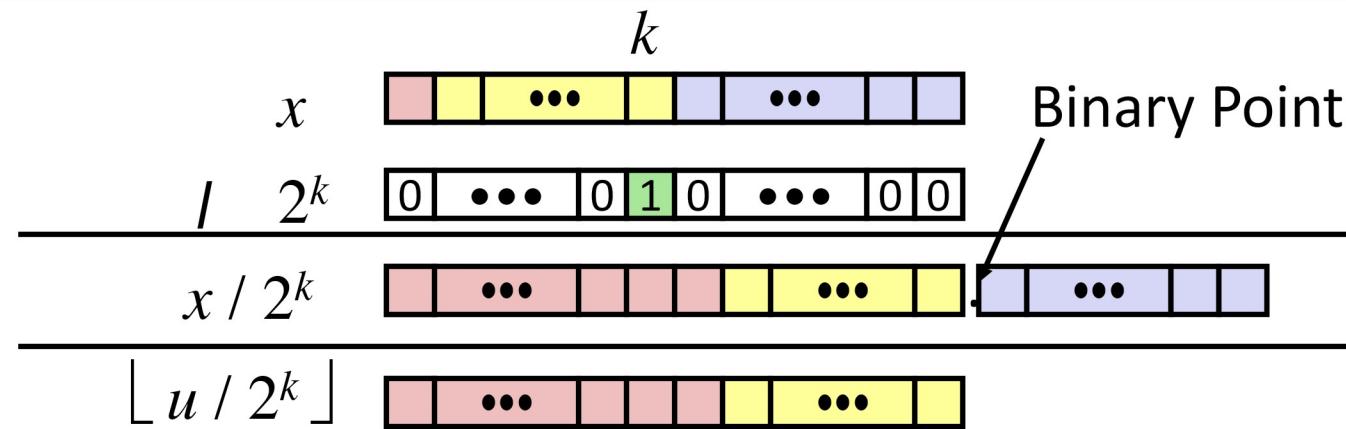
```
movq %rdi, %rax  
shrq $3, %rax
```

## Explanation

```
# Logical shift  
return x >> 3;
```

# Signed power-of-2 divide with shift

Operands:



## ■ Quotient of signed by power-of-2

- $x \gg k$  gives  $\lfloor x / 2^k \rfloor$

## ■ Uses **arithmetic** shift

# Signed power-of-2 divide with shift

	<b>Division</b>	<b>Computed</b>	<b>Hex</b>	<b>Binary</b>
<b>y</b>	-15 213	-15 213	C4 93	11000100 10010011
<b>y &gt;&gt; 1</b>	-7 606.5	-7 607	E2 49	11100010 01001001
<b>y &gt;&gt; 4</b>	-950.8125	-951	FC 49	11111100 01001001
<b>y &gt;&gt; 8</b>	-59.4257813	-60	FF C4	11111111 11000100

- For a negative number, arithmetic right shift rounds down rather than toward zero
  - For a positive number, we have 0 as the most significant bit, and so the effect is the same as for a logical right shift

# Correct power-of-2 divide with shift

- We can correct for this improper rounding by “biasing” the value before shifting
- Quotient of negative number by power-of-2
  - Want  $\lceil x / 2^k \rceil$  (Round toward 0 – convention for integer division)
  - Compute as  $\lfloor (x+2^k-1) / 2^k \rfloor$ 
    - In C: `(x + (1<<k)-1) >> k`
- This technique exploits the property that  $\lceil x/y \rceil = \lfloor (x + y - 1)/y \rfloor$  for integers  $x$  and  $y$  such that  $y > 0$

# Correct power-of-2 divide with shift

Case 1: No rounding is required ( $x \geq 0$ )

Dividend:

$$\begin{array}{r}
 u \quad \begin{array}{|c|c|c|c|c|c|c|c|} \hline & 1 & \cdots & 0 & \cdots & 0 & 0 \\ \hline \end{array} \\
 +2^k - 1 \quad \begin{array}{|c|c|c|c|c|c|c|c|} \hline & 0 & \cdots & 0 & 0 & 1 & \cdots & 1 & 1 \\ \hline \end{array} \\
 \hline
 \begin{array}{|c|c|c|c|c|c|c|c|} \hline & 1 & \cdots & 1 & \cdots & 1 & 1 \\ \hline \end{array} \quad \text{Binary Point}
 \end{array}$$

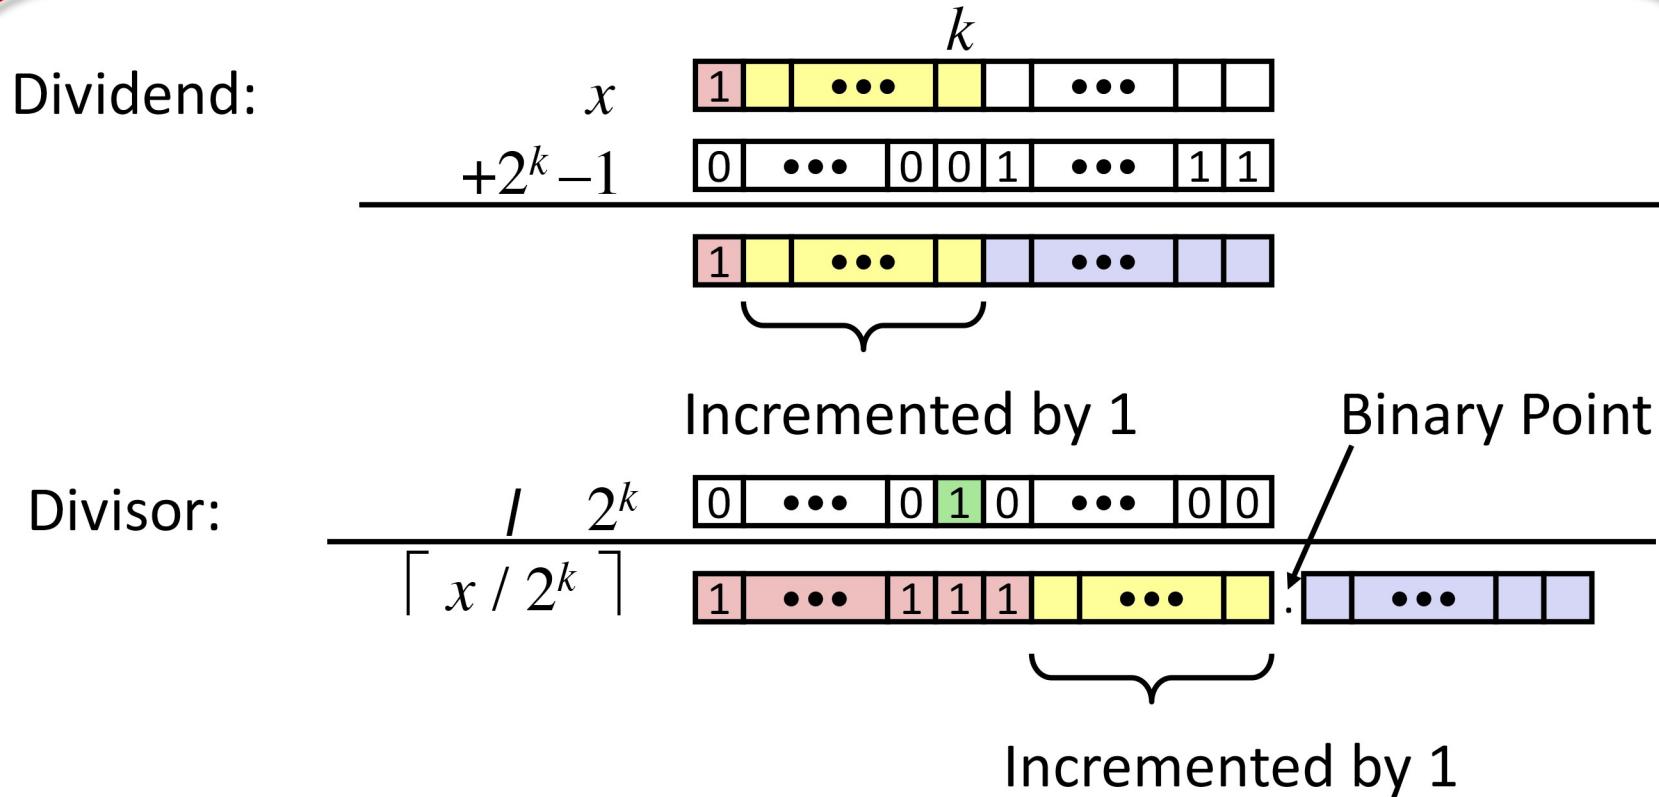
Divisor:

$$\begin{array}{r}
 / \quad 2^k \quad \begin{array}{|c|c|c|c|c|c|c|c|} \hline & 0 & \cdots & 0 & 1 & 0 & \cdots & 0 & 0 \\ \hline \end{array} \\
 \hline
 \lceil u / 2^k \rceil \quad \begin{array}{|c|c|c|c|c|c|c|c|} \hline & 1 & \cdots & 1 & 1 & 1 & \cdots & 1 & 1 \\ \hline \end{array} . \quad \begin{array}{|c|c|c|c|c|c|c|c|} \hline & 1 & \cdots & 1 & 1 \\ \hline \end{array}
 \end{array}$$

*Biasing has no effect*

# Correct power-of-2 divide with shift

Case 2: Rounding is required ( $x < 0$ )



*Biasing adds 1 to final result*

# Compiled signed division code

## C function

```
long idiv8(long x) {  
    return x/8;  
}
```

## Compiled arithmetic operations

```
movq %rdi, %rax  
testl %rax, %rax  
js L4  
L3:  
    sarq $3, %rax  
    ret  
L4:  
    addq $7, %rax  
    jmp L3
```

## Explanation

```
if x < 0  
    x += 7;  
# Arithmetic shift  
return x >> 3;
```

# Arithmetic: Basic rules

## ■ Addition

- **Unsigned/signed:** Normal addition followed by truncate,  
same operation on bit level
- **Unsigned:** addition mod  $2^w$ 
  - Mathematical addition + possible subtraction of  $2^w$
- **Signed:** modified addition mod  $2^w$  (result in proper range)
  - Mathematical addition + possible addition or subtraction of  $2^w$

## ■ CPUs can use the same machine instruction to perform either unsigned or signed addition

# Arithmetic: Basic rules

## ■ Multiplication

- **Unsigned/signed**: Normal multiplication followed by truncate, same operation on bit level
- **Unsigned**: multiplication mod  $2^w$
- **Signed**: modified multiplication mod  $2^w$  (result in proper range)

## ■ On most machines, integer multiplication and division are fairly slow

- Replace them by constant factors with combinations of shift and addition/subtraction operations

# Arithmetic: Basic rules

## ■ Left shift

- Unsigned/signed: multiplication by  $2^k$
- Always logical shift

## ■ Right shift

- Unsigned: logical shift, div (division + round to zero) by  $2^k$
- Signed:
  - Positive numbers: div (division + round to zero) by  $2^k$
  - Negative numbers: div (division + round away from zero) by  $2^k$ 
    - Use biasing to fix