

Floating Point

Arquitectura de Computadores
Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto

Luís Nogueira (lmn@isep.ipp.pt)

Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary
- Practice problem

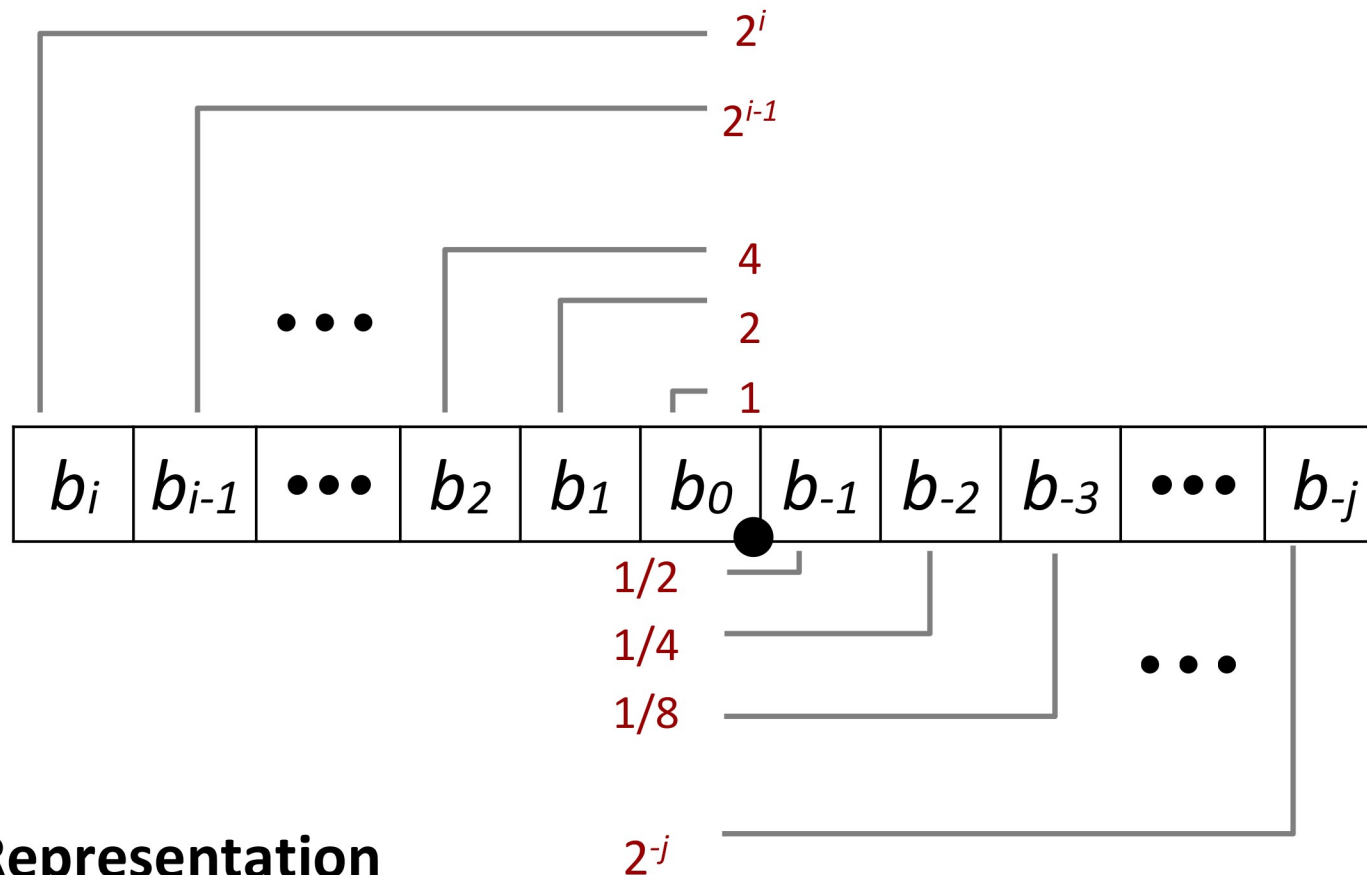
Introduction

- **A floating-point representation encodes rational numbers of the form $V = x \times 2^y$**
- **It is useful for performing computations involving**
 - very large numbers ($|V| \gg 0$)
 - numbers very close to 0 ($|V| \ll 1$)
 - and more generally as an approximation to real arithmetic
- **What issues come up when trying to devise a data representation for floating point numbers?**
 - It turns out these issues are more complicated than representing integers

Fractional binary numbers

- A first step in understanding floating-point numbers is to consider binary numbers having fractional values
- What is 1011.101_2 ?

Fractional binary numbers



■ Representation

- Bits to right of “binary point” represent fractional powers of 2

- Represents rational number: $\sum_{k=-j}^i b_k \times 2^k$

Fractional binary numbers

- A first step in understanding floating-point numbers is to consider binary numbers having fractional values
- What is 1011.101_2 ?
- The binary representation of 11.625

$$1 \times 2^3 + 1 \times 2^1 + 1 \times 2^0 + 1 \times 2^{-1} + 1 \times 2^{-3} =$$

$$8 + 2 + 1 + 1/2 + 1/8 =$$

$$11.625$$

Fractional binary numbers: Examples

■ Value Representation

5.75 = 5 3/4	101.11 ₂
2.875 = 2 7/8	10.111 ₂
1.4375 = 1 7/16	1.0111 ₂

■ Observations

- Divide by 2 by shifting binary point left (unsigned)
- Multiply by 2 by shifting binary point right
- Numbers of form 0.111111...₂ are just below 1.0
 - $1/2 + 1/4 + 1/8 + \dots + 1/2^i + \dots \rightarrow 1.0$
 - Use notation $1.0 - \epsilon$

Representable numbers

■ Limitation #1

- Can only exactly represent numbers of the form $x/2^k$
 - Other rational numbers have repeating bit representations

Value	Representation
1/3	0.0101010101 [01] ... ₂
1/5	0.001100110011 [0011] ... ₂
1/10	0.0001100110011 [0011] ... ₂

■ Limitation #2

- Just one setting of binary point within the w bits
 - Not efficient for representing very large or very small numbers

Today: Floating Point

- Background: Fractional binary numbers
- **IEEE floating point standard: Definition**
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary
- Practice problem

IEEE Floating Point

■ IEEE Standard 754

- Established in 1985 as uniform standard for floating point arithmetic
 - Before that, many idiosyncratic formats
- Supported by all major CPUs

■ Driven by numerical concerns

- Nice standards for rounding, overflow, underflow
- Hard to make fast in hardware
 - Numerical analysts predominated over hardware designers in defining standard

Floating point representation

■ Numerical form:

$$(-1)^s M 2^E$$

- **Sign bit** s determines whether number is negative ($s = 1$) or positive ($s = 0$)
- **Significand** M normally a fractional value in range $[1.0, 2.0)$.
- **Exponent** E weights value by power of two

■ Encoding

- MSB s is sign bit s
- **exp** field encodes E (but is not equal to E)
- **frac** field encodes M (but is not equal to M)



Precision options

■ Single precision: 32 bits (float type in C)



■ Double precision: 64 bits (double type in C)



■ Extended precision: 80 bits (Intel only)



Normalized values

$$v = (-1)^s M 2^E$$

- The common case, when $exp \neq 000\dots 0$ and $exp \neq 111\dots 1$

- Exponent coded as a *biased* value: $E = Exp - Bias$

- Exp : unsigned value of exp field
- $Bias = 2^{k-1} - 1$, where k is number of exponent bits
 - Single precision: 127 (Exp : 1...254, E : -126...127)
 - Double precision: 1023 (Exp : 1...2046, E : -1022...1023)

- Significand coded with implied leading 1: $M = 1.xxx\dots x_2$

- $xxx\dots x$: bits of $frac$ field
- Minimum when $frac = 000\dots 0$ ($M = 1.0$)
- Maximum when $frac = 111\dots 1$ ($M = 2.0 - \epsilon$)
- Get extra leading bit for “free”

Normalized encoding example

$$v = (-1)^s M 2^E$$

$$E = \text{Exp} - \text{Bias}$$

■ **Value:** `float f = 15213.0;`

$$15213_{10} = 11101101101101_2$$

$$= 1.1101101101101_2 \times 2^{13}$$

■ **Significand**

$$M = 1.\underline{1101101101101}_2$$

$$\text{frac} = \underline{1101101101101}0000000000_2 \quad (23 \text{ bits})$$


■ **Exponent**

$$E = 13$$

$$\text{Bias} = 127$$

$$\text{Exp} = 140 = 10001100_2 \quad (8 \text{ bits})$$

■ **Result:**



Denormalized values

$$v = (-1)^s M 2^E$$

$$E = 1 - Bias$$

- Condition: $exp = 000...0$
- Exponent value: $E = 1 - Bias$ (instead of $E = 0 - Bias$)
- Significand coded with implied leading 0: $M = 0.xxx...x_2$
 - $xxx...x$: bits of $frac$
- Case: $exp = 000...0, frac = 000...0$
 - Represents zero value
 - Note distinct values: +0 and -0 (why?)
- Case: $exp = 000...0, frac \neq 000...0$
 - Numbers closest to 0.0
 - Equispaced

Denormalized values

$$v = (-1)^s M 2^E$$

$$E = 1 - \text{Bias}$$

- **Condition: $\text{exp} = 111\dots 1$**

- **Case: $\text{exp} = 111\dots 1, \text{frac} = 000\dots 0$**
 - Represents value ∞ (infinity)
 - Operation that overflows
 - Both positive and negative
 - E.g., $1.0/0.0 = -1.0/-0.0 = +\infty$, $1.0/-0.0 = -\infty$

- **Case: $\text{exp} = 111\dots 1, \text{frac} \neq 000\dots 0$**
 - Not-a-Number (NaN)
 - Represents case when no numeric value can be determined
 - E.g., $\text{sqrt}(-1)$, $\infty - \infty$, $\infty \times 0$

Visualization: Floating point encodings


■ Normalized



■ Denormalized




■ Infinity



■ NaN




Visualization: Floating point encodings



Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- **Example and properties**
- Rounding, addition, multiplication
- Floating point in C
- Summary
- Practice problem

Tiny floating point example



■ 8-bit floating point representation

- The sign bit is in the most significant bit
- The next four bits are the exponent, with a bias of $2^4 - 1 = 7$
- The last three bits are the **frac**

■ Same general form as IEEE Format

- Normalized, denormalized
- Representation of 0, NaN, infinity

Dynamic range (positive only)

$$v = (-1)^s M 2^E$$

n: $E = \text{Exp} - \text{Bias}$
d: $E = 1 - \text{Bias}$

closest to zero

largest denormalized
smallest normalized

closest to 1 below

closest to 1 above

largest normalized

Denormalized
numbers

Normalized
numbers

s	exp	frac	E	Value
0	0000	000	-6	0
0	0000	001	-6	$1/8 * 1/64 = 1/512$
0	0000	010	-6	$2/8 * 1/64 = 2/512$
...				
0	0000	110	-6	$6/8 * 1/64 = 6/512$
0	0000	111	-6	$7/8 * 1/64 = 7/512$
0	0001	000	-6	$8/8 * 1/64 = 8/512$
0	0001	001	-6	$9/8 * 1/64 = 9/512$
...				
0	0110	110	-1	$14/8 * 1/2 = 14/16$
0	0110	111	-1	$15/8 * 1/2 = 15/16$
0	0111	000	0	$8/8 * 1 = 1$
0	0111	001	0	$9/8 * 1 = 9/8$
0	0111	010	0	$10/8 * 1 = 10/8$
...				
0	1110	110	7	$14/8 * 128 = 224$
0	1110	111	7	$15/8 * 128 = 240$
0	1111	000	n/a	inf

Interesting numbers

{single, double}

<i>Description</i>	<i>exp</i>	<i>frac</i>	<i>Numeric Value</i>
■ Zero	00...00	00...00	0.0
■ Smallest positive denorm.	00...00	00...01	$2^{-\{23,52\}} \times 2^{-\{126,1022\}}$
<ul style="list-style-type: none"> ■ Single $\approx 1.4 \times 10^{-45}$ ■ Double $\approx 4.9 \times 10^{-324}$ 			
■ Largest denormalized	00...00	11...11	$(1.0 - \epsilon) \times 2^{-\{126,1022\}}$
<ul style="list-style-type: none"> ■ Single $\approx 1.18 \times 10^{-38}$ ■ Double $\approx 2.2 \times 10^{-308}$ 			
■ Smallest positive normalized	00...01	00...00	$1.0 \times 2^{-\{126,1022\}}$
<ul style="list-style-type: none"> ■ Just larger than largest denormalized 			
■ One	01...11	00...00	1.0
■ Largest Normalized	11...10	11...11	$(2.0 - \epsilon) \times 2^{\{127,1023\}}$
<ul style="list-style-type: none"> ■ Single $\approx 3.4 \times 10^{38}$ ■ Double $\approx 1.8 \times 10^{308}$ 			

Special properties of the IEEE encoding

■ Floating-point zero same as integer zero

- All bits are equal to 0

■ Can (almost) use unsigned integer comparison

- Must first compare sign bits
- Must consider $-0 = 0$
- NaNs problematic
 - Will be greater than any other values, including infinity
 - What should comparison yield?
- Otherwise OK
 - Denormalized vs. normalized
 - Normalized vs. infinity

Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- **Rounding, addition, multiplication**
- Floating point in C
- Summary

Floating point operations: Basic idea

$$\blacksquare \mathbf{x} +_{\mathbf{f}} \mathbf{y} = \text{Round}(\mathbf{x} + \mathbf{y})$$

$$\blacksquare \mathbf{x} \times_{\mathbf{f}} \mathbf{y} = \text{Round}(\mathbf{x} \times \mathbf{y})$$

■ Basic idea

- First **compute exact result**
- Make it fit into desired precision
 - Possibly overflow if exponent too large
 - Possibly **round to fit into *frac***

Rounding

- For a value x , find the “closest” matching value x' that can be represented in the desired floating-point format
 - One key problem is to define the direction to round a value that is halfway between two possibilities

- IEEE floating-point rounding modes

	1.40	1.60	1.50	2.50	−1.50
■ Towards zero	1	1	1	2	−1
■ Round down ($-\infty$)	1	1	1	2	−2
■ Round up ($+\infty$)	2	2	2	3	−1
■ Round to nearest, ties to even (default)	1	2	2	2	−2

Closer look at round-to-even

■ Default rounding mode

- Hard to get any other kind without dropping into assembly
- All others are statistically biased
 - Sum of set of positive numbers will consistently be over- or under-estimated

■ Applying to other decimal places / bit positions

- When exactly halfway between two possible values
 - Round so that least significant digit is even
- E.g., round to nearest hundredth

7.8949999	7.89	(Less than half way)
7.8950001	7.90	(Greater than half way)
7.8950000	7.90	(Half way—round up)
7.8850000	7.88	(Half way—round down)

Rounding binary numbers

■ Binary fractional numbers

- “Even” when least significant bit is 0
- “Half way” when bits to right of rounding position = 100...₂

■ Examples

- Round to nearest 1/4 (2 bits right of binary point)

Value	Binary	Rounded	Action	Rounded Value
2 3/32	10.00011 ₂	10.00 ₂	(<1/2—down)	2
2 3/16	10.00110 ₂	10.01 ₂	(>1/2—up)	2 1/4
2 7/8	10.11100 ₂	11.00 ₂	(1/2—up)	3
2 5/8	10.10100 ₂	10.10 ₂	(1/2—down)	2 1/2

Floating point multiplication

■ $(-1)^{s1} M1 2^{E1} \times (-1)^{s2} M2 2^{E2}$

■ **Exact Result:** $(-1)^s M 2^E$

- Sign s : $s1 \wedge s2$
- Significand M : $M1 \times M2$
- Exponent E : $E1 + E2$

■ Fixing

- If $M \geq 2$, normalize - shift M right, increment E
- If E out of range, overflow
- Round M to fit **frac** precision

■ Implementation

- Biggest inconvenience is multiplying significands

Floating point addition


$$\blacksquare (-1)^{s1} M1 2^{E1} + (-1)^{s2} M2 2^{E2}$$

- Assume $E1 > E2$

$$\blacksquare \text{Exact Result: } (-1)^s M 2^E$$

- Sign s , significand M :
 - Result of signed align & add
- Exponent E : $E1$

Get binary points lined up



Fixing

- If $M \geq 2$, normalize - shift M right, increment E
- if $M < 1$, normalize - shift M left k positions, decrement E by k
- Overflow if E out of range
- Round M to fit **frac** precision

Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- **Floating point in C**
- Summary

Floating Point in C

■ C guarantees two levels

- `float` single precision
- `double` double precision

■ Conversions/casting

- Casting between `int`, `float`, and `double` changes bit representation
- `double/float → int`
 - Truncates fractional part
 - Like rounding toward zero
 - Not defined when out of range or NaN: Generally sets to TMin
- `int → double`
 - Exact conversion, as long as `int` has ≤ 53 bit word size
- `int → float`
 - Will round according to rounding mode

Floating Point in C: Exercise

■ For each of the following C expressions, either:

- Argue that it is true for all argument values
- Explain why not true

```
int x = ...;
float f = ...;
double d = ...;
```

Assume neither
d nor **f** is NaN

- ☐ `x == (int)(float) x`
- ☐ `x == (int)(double) x`
- ☐ `f == (float)(double) f`
- ☐ `d == (double)(float) d`
- ☐ `f == -(-f);`
- ☐ `2/3 == 2/3.0`
- ☐ `d < 0.0 ⇒ ((d*2) < 0.0)`
- ☐ `d > f ⇒ -f > -d`
- ☐ `d * d >= 0.0`
- ☐ `(d+f) - d == f`

Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- **Summary**
- Practice problem


Summary

- IEEE Floating Point has clear mathematical properties
- Represents numbers of form $(-1)^S M 2^E$
- One can reason about operations independent of implementation
 - As if computed with perfect precision and then rounded
- Not the same as real arithmetic
 - Violates associativity/distributivity
 - Makes life difficult for compilers & serious numerical applications programmers

Today: Floating Point

- Background: Fractional binary numbers
- IEEE floating point standard: Definition
- Example and properties
- Rounding, addition, multiplication
- Floating point in C
- Summary
- **Practice problem**


Practice problem



■ Convert 8-bit unsigned numbers to tiny floating point format

<i>Value</i>	<i>Binary</i>
128	10000000
15	00001101
33	00010001
35	00010011
138	10001010
63	00111111


Practice problem



■ Steps:

1. Normalize to have leading 1
2. Round to fit within fraction
3. Postnormalize to deal with effects of rounding

Step 1: Normalize




■ Requirement

- Set binary point so that numbers of form 1.xxxxx
- Adjust all to have leading one
 - Decrement exponent as shift left

<i>Value</i>	<i>Binary</i>	<i>Fraction</i>	<i>Exponent</i>
128	10000000	1.0000000	7
15	00001101	1.1010000	3
17	00010001	1.0001000	4
19	00010011	1.0011000	4
138	10001010	1.0001010	7
63	00111111	1.1111100	5

Step 2: Rounding



■ Round up conditions

- Round = 1, Sticky = 1 \rightarrow > 0.5
- Guard = 1, Round = 1, Sticky = 0 \rightarrow Round to even

<i>Value</i>	<i>Fraction</i>	<i>GRS</i>	<i>Incr?</i>	<i>Rounded</i>
128	1.000 0000	000	N	1.000
15	1.101 0000	100	N	1.101
17	1.000 1000	010	N	1.000
19	1.001 1000	110	Y	1.010
138	1.000 1010	011	Y	1.001
63	1.111 1100	111	Y	10.000

Step 3: Postnormalize

■ Issue

- Rounding may have caused overflow
- Handle by shifting right once & incrementing exponent

<i>Value</i>	<i>Rounded</i>	<i>Exp</i>	<i>Adjusted</i>	<i>Result</i>
128	1.000	7		128
15	1.101	3		15
17	1.000	4		16
19	1.010	4		20
138	1.001	7		134
63	10.000	5	1.000/6	64