

ARQCP Course

Arquitetura de Computadores
Licenciatura em Engenharia Informática

2023/24
Paulo Baltarejo Sousa
`pbs@isep.ipp.pt`

ISEP INSTITUTO SUPERIOR
DE ENGENHARIA DO PORTO

Material and Slides

Some of the material/slides are adapted from various:

- Presentations found on the internet;
- Books;
- Web sites;
- ...

- 1 Definitions and Concepts
- 2 Computer program
- 3 Central Processing Unit (CPU)
- 4 Why Assembly?
- 5 Coding in Assembly

Definitions and Concepts

- RISC (**Reduced Instruction Set Computer**) and CISC (**Complex Instruction Set Computer**) refer to the processor architectures that utilize different data processing instruction sets to perform basic logical and input/output operations.
- CISC
 - This architecture was introduced in the 1970s by Intel Corporation when the earliest computers focused on enhancing CPU speed by **minimizing the number of instructions** per program.
 - It can perform **multiple operations in a single instruction**.
- RISC
 - These were introduced in the 1980s by David Patterson and John Hennessy to overcome the complexities of CISC processors.
 - RISC processors work with more instructions (simple).
 - It performs **a single operation in an instruction**.

- An Instruction Set Architecture (ISA) **is part of the abstract model of a computer that defines how the CPU is controlled by the software.**
- The ISA acts as an **interface between the hardware and the software, specifying both what the processor is capable of doing as well as how it gets done.**
- Example ISAs:
 - Intel: x86, IA32, Itanium, x86-64
 - ARM: used in almost all mobile phones, Apple M1
 - RISC V: new open-source ISA
 - MIPS

- **x86 is a type of ISA** for computer processors originally developed by Intel in 1978.
- The x86 architecture is based on Intel's 8086 (hence the name) microprocessor.
- At first, it was a 16-bit instruction set for 16-bit processors, and later it grew to 32-bit instruction sets.
- The number of bits signifies how much information the CPU can process per cycle.
 - For example, a 32-bit CPU transfers up to 32 bits of data per clock cycle.
- The x86 architecture's most significant limitation is that it can handle a maximum of 4GB of RAM.

- **x64 (short for x86-64) is an instruction set architecture based on x86, extended to enable 64-bit code.**
- It was first released in 2000, introducing two modes of operation - the 64-bit mode and the compatibility mode, which allows users to run 16-bit and 32-bit applications as well.
 - Since the entire x86 instruction set remains implemented in the x64 one, the older executables run with practically no performance penalties.
- The x64 architecture supports much greater amounts of virtual and physical memory than the x86 architecture, allowing applications to store large data amounts in memory.
- Additionally, x64 expands the number of general-purpose registers to 16, providing further enhancements and functionality.
- The x64 architecture allows the CPU to process 64-bits of data per clock cycle.

x86 vs. x86-64 Architectures



Supports up to 4 GB of RAM



Supports up to 16 EB of RAM

Has a 32-bit bus for transmitting up to 32 bits of data in a single cycle



Has a 64-bit bus allowing parallel transmits of large data amounts

Supports only 32-bit software



Supports 32-bit and 64-bit software


Slower processing speeds



Fast processing speeds

- A microarchitecture **is the digital logic that allows an instruction set to be executed.**
- It is the combined **implementation of registers, memory, arithmetic logic units, multiplexers, and any other digital logic blocks.**
 - All of this, together, **forms the processor.**
- A microarchitecture **combined with an instruction set architecture (ISA) makes up the system's computer architecture as a whole.**
- Different microarchitectures **can implement the same ISA, but with trade-offs in things like power efficiency or execution speed.**

- Pipelining is a technique used to **improve the execution throughput of a CPU by using the processor resources in a more efficient manner.**
- The basic idea is to **split the processor instructions into a series of small independent stages.**
- Each stage is designed to perform a certain part of the instruction.



- **Fetch** an instruction from memory
- **Decode** the instruction to be executed
- **Execute** the instruction
- **Write** the result back to register or memory


- Branch prediction is a technique used in CPU design that attempts **to guess the outcome of a conditional operation and prepare for the most likely result**.
- A digital circuit that performs this operation is known as a **branch predictor**.
- How does it work?
 - When a conditional operation such as an `if...else` statement needs to be processed, the branch predictor "speculates" which condition most likely will be met.
 - It then executes the operations required by the most likely result ahead of time.
 - This way, they are already complete if and when the guess was correct.
 - At runtime, if the guess turns out not to be correct, the CPU executes the other branch of operation, incurring a slight delay.
 - But if the guess was correct, speed is significantly increased.

- A processor that executes the instructions one after the other, may use the resources inefficiently that leads to poor performance of the processor.
- Out-of-order execution can improve the performance of the processor.
- Out-of-order execution can be achieved by executing the instruction in an different from the original order they appear.
- This approach efficiently uses instruction cycles (fetch, decode, execute and write) and reduces costly delay.
- A processor will execute the instructions in an order of availability of data or operands instead of original order of the instructions in the program.
- By doing so the processor will avoid being idle while data is retrieved for the next instruction in a program.

Computer program


Execution

- A computer **program** is a collection of instructions (and data) that performs a specific task when executed by a computer.
- The **program instructions** are loaded into main memory (RAM).
- The CPU must look into **memory and fetch in the instructions and data and act upon them according to what the instructions are.**



Compilation (I)

- The gcc C compiler generates its output in the form of **assembly code**, a textual representation of the machine code giving the individual instructions in the program.
- gcc then invokes both an **assembler** and a **linker** to generate the **executable machine code** from the assembly code.



- **Machine language**, or **machine code**, is the most basic set of instructions that a computer can execute.
- **Each type of processor** has its own set of **machine language instructions** (ISA).

C

```
#include<stdio.h>
int main(){
    printf("Hello world!\n");
    return 0;
}
```

Machine code

```
01111111 01000101 01001100 01000110 00000010 00000001
00000001 00000000 00000000 00000000 00000000 00000000
00000000 00000000 00000000 00000000 00000001 00000000
...
```

Assembly


```
.section .data
.LC0:
.string "Hello world!"
.section .text
.global main
main:
    pushl %ebp
    movl %esp, %ebp
    movl $.LC0, %edi
    call puts
    movl $0, %eax
    popl %ebp
    ret
```

- Assembly language **is abstraction layer on machine code instructions.**
- An assembly language program is **stored as text file** (just as a higher level language program).
- Each **assembly instruction represents exactly one machine instruction** (according to the ISA).
- The general form of an assembly instruction is:

mnemonic operands


- For example, an addition operation could be coded in assembly language as:
- addl %eax, %ebx
- The word `addl` is a **mnemonic** for the addition instruction.
 - An assembler **is a program that reads a text file with assembly instructions and converts the assembly into machine code.**

The general format of machine code instruction:

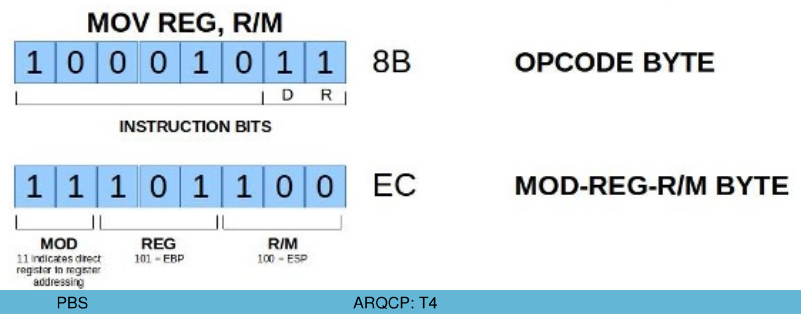


- An **OpCode** (abbreviated from operation code) is the portion of a machine language instruction that specifies the operation to be performed.
- The **MOD-R/M** byte specifies instruction operands and their addressing mode.
- ...

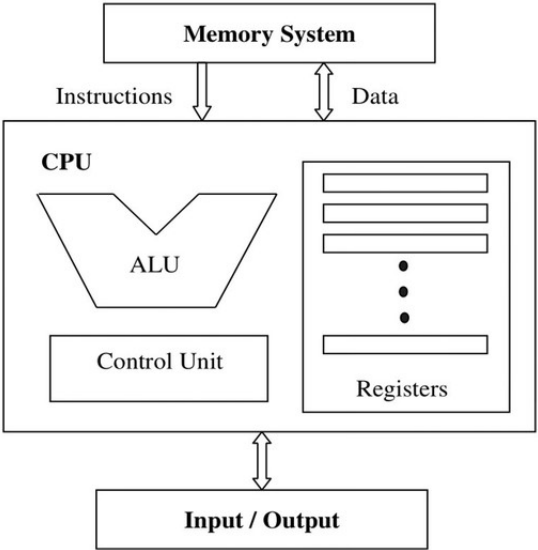
Instruction to disassemble: 55 → PUSH EBP



Instruction to disassemble: 8B EC → MOV EBP, ESP




Central Processing Unit (CPU)



- The Control Unit is the part of the computer's CPU, which **directs the operation of the processor**.
- It coordinates the flow of data out of, into, and between the various (sub)units of a processor.
- It **understands commands and instructions**.
- It **regulates the flow of data within the processor**.
- It accepts external commands or instructions, which it turns into a series of control signals.
- It is in charge of a CPU's multiple execution units (such as ALUs, data buffers, and registers).
- It also performs a variety of activities, including fetching, decoding, handling execution, and storing results.

- The Arithmetic and Logic Unit is a unit of central processing unit where all arithmetic and logical operations are carried out.

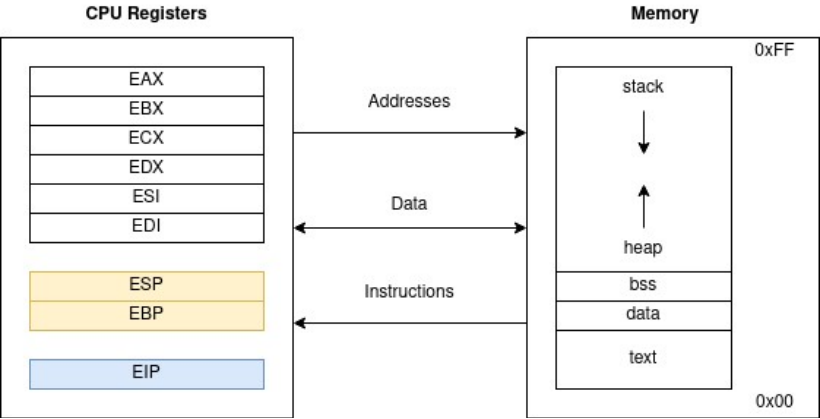


- The basic ALU operates on integer numbers.
 - Two numbers (operands), A and B, are presented to the input of the ALU, and also an instruction - formally called an **opcode**, such as "add", "subtract", ...
 - The ALU carries out the requested operation on A and B and produces an output result.
 - Along with the result, a set of status flags are set which include 'Zero', 'Carry', 'Negative', 'Overflow'.

- The **x86 architecture contains fourteen registers.**
- **Data Registers:** Holds data for operations.
 - Accumulator Register (AX), Base Register (BX), Count Register (CX), and Data Register (DX).
- **Address Register:** Holds Address of instruction.
 - Segment Registers
 - Code Segment (CS), Data Segment (DS), Stack Segment (SS), and Extra Segment (ES, FS, GS)
 - Pointer Registers
 - Stack Pointer (SP), Base Pointer (BP), and Instruction Pointer (IP)
 - Index Registers
 - Source Index (SI) and Destination Index (DI)
- **Status Registers:** Store the current status of the processor.
 - Status Flags
 - Control Flags


Data			Address		Status	
EAX	AX		EBP	BP	EFLAGS	FLAGS
	AH	AL				
EBX	BX		ESP	SP		
	BH	BL				
ECX	CX		EIP	IP		
	CH	CL				
EDX	DX			CS		
	DH	DL		DS		
ESI	SI			SS		
EDI	DI			ES		
				FS		
				GS		

- CPU and Memory **are the core** of computational system.



- Multipurpose Registers (EAX, EBX, ..., EDI)
- Special-purpose Registers (ESP, EBP, and EIP)

- The **Instruction Pointer (EIP)** register contains the **memory address of the next instruction to be executed**




- **After the execution of the instruction, the EIP register is automatically increased to the address of next instruction, unless it ... jumps**


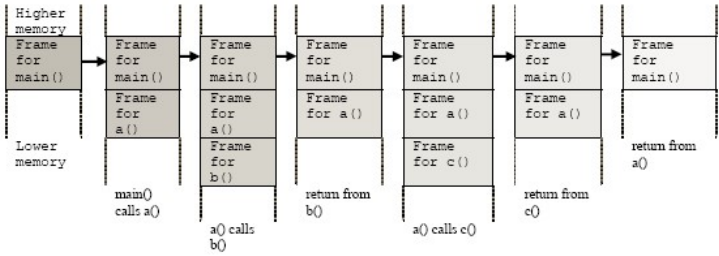
```
movl $1, %eax
jmp label
movl $3, %ebx
label:
int $0x80
```

- **Stack Pointer (ESP)** register contains the **address of the top of the stack.**
- **Base Pointer (EBP)** register contains the **address of the bottom of the stack frame**

```
last() {  
}  
two() {  
    last();  
}  
one() {  
    two();  
}  
main() {  
    one();  
}
```



```
int c() {
    return 0;
}
int b() {
    return 0;
}
int a() {
    b();
    c();
    return 0;
}
int main() {
    a();
    return 0;
}
```



```
b:
    pushl %ebp
    movl %esp,%ebp
    movl $0, %eax
    movl %ebp,%esp
    popl %ebp
    ret
a:
    pushl %ebp
    movl %esp,%ebp
    ...
    call b
    ...
    movl %ebp,%esp
    popl %ebp
    ret
```

- The **Return address** is the address of the next instruction right after `call`
- `call <label>`
 - Push **Return address** on stack.
 - Changes the EIP register with the address represented by `label`
- `ret` instruction **pops the address from stack and set EIP register with popped value.**
 - It should be the **Return address** previously pushed by `call` instruction.

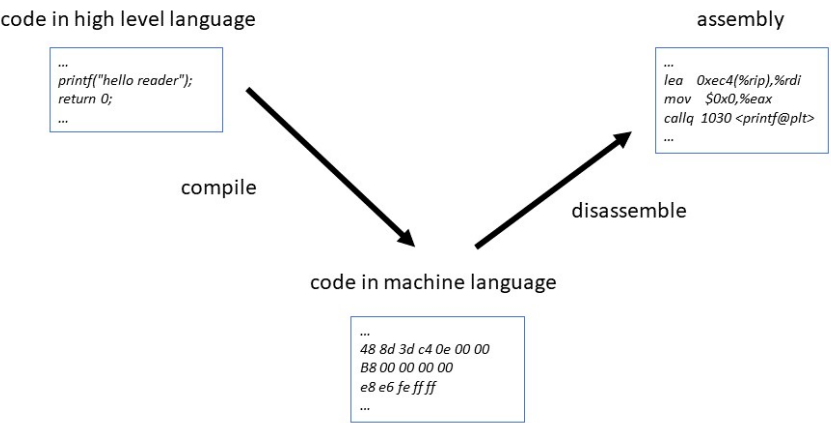
Why Assembly?

Why is learning assembly language still important?

- Assembly machine code **is the foundation of all software**, the fundamental language that computer hardware understands.
 - As programmers and developers, it's essential to grasp how it works.
- The need for programmers **to learn assembly code has shifted over the years** from one of being able to **write programs** directly in assembly to one of being able to read and understand the **code generated by compilers**.
- Understand the **compiler optimizations**.
- Better understanding of architecture issues.
 - Understand how **computer systems perform**.
 - Understanding processor and memory function.
- Improve **algorithm development** skills.
- And so on.

- Reverse engineering is the **inverse process to the normal development of a product**, i.e. the process **begins with a finished product and it is dismantled into its individual parts**.
- Reverse engineering refers **to taking an already compiled code and converting it back into a human-readable format**.
 - In most cases, reverse engineering allows us to understand the program's functionality better and determine how it runs.
- One of the programming languages that are used in reverse engineering is **Assembly Language**

- Disassembler
 - Converting the program **from machine code to assembly code** in order to analyze the program.



■ Disassemble a Binary File: objdump command

■ > gcc hello.c -o hello

```
#include<stdio.h>
int main(){
    printf("Hello world!\n");
    return 0;
}
```

■ > objdump -d hello

```
...
0000000000001149 <main>:
1149: f3 0f 1e fa  endbr64
114d: 55  push %rbp
114e: 48 89 e5  mov %rsp,%rbp
1151: 48 8d 3d ac 0e 00 00  lea 0xaeac(%rip),%rdi # 2004 <_IO_stdin_used+0x4>
1158: e8 f3 fe ff ff  callq 1050 <puts@plt>
115d: b8 00 00 00 00  mov $0x0,%eax
1162: 5d  pop %rbp
1163: c3  retq
1164: 66 2e 0f 1f 84 00 00  nopw %cs:0x0(%rax,%rax,1)
116b: 00 00 00
116e: 66 90  xchg %ax,%ax
...
```

Coding in Assembly

- Assembly language is a **low-level programming language that closely corresponds to the machine code instructions of a specific CPU type.**
- It acts as an intermediary layer between high-level languages and machine code, providing a more human-readable representation of the machine code instructions.
- Each assembly language **is designed for exactly one specific computer architecture.**
- Unlike high-level languages, **assembly language is not portable.**
- A program written **in assembly language can only be run on the same type of machine that it was written for.**

- Writing in assembly language involves using **mnemonic codes to represent each low-level machine operation**.
- These mnemonics **are then translated directly into machine code** by an assembler.
- The syntax of assembly language consists of instructions that are made up of an operation code (opcode) followed by a list of arguments.
 - The **opcode**, or mnemonic, is an abbreviated name for an operation (e.g., ADD for addition, MOV for moving data).
 - The arguments are typically **memory addresses, constants, or CPU registers**.
- Assembly language allows programmers to use symbolic addresses (**labels**), which means that they don't need to know the exact memory location they are coding for.
 - This makes the code more readable and easier to debug and maintain.
- All of the lines beginning with "." are directives to guide the assembler and linker.

- An assembly program can be divided into three sections:

- The `data` section

- It is used for declaring initialized data or constants.
- This data does not change at runtime.
- It is used for declaring constant values, file names or buffer size etc.

- The `bss` section

- It is used for declaring uninitialized variables

- The `text` section

- It is used for code (instructions and logic)

