

ARQCP Course

Arquitetura de Computadores
Licenciatura em Engenharia Informática

2023/24
Paulo Baltarejo Sousa
`pbs@isep.ipp.pt`

ISEP INSTITUTO SUPERIOR
DE ENGENHARIA DO PORTO

Material and Slides

Some of the material/slides are adapted from various:

- Presentations found on the internet;
- Books;
- Web sites;
- ...

1 x86-64

2 The 64 bit x86 C Calling Convention

x86-64

Registers (I)


64	32	16	8	0	64	32	16	8	0
rax	eax	ax	ah	al	r8	r8d	r8w	r8b	
rbx	ebx	bx	bh	bl	r9	r9d	r9w	r9b	
rcx	ecx	cx	ch	cl	r10	r10d	r10w	r10b	
rdx	edx	dx	dh	dl	r11	r11d	r11w	r11b	
rsi	esi	si	sil		r12	r12d	r12w	r12b	
rdi	edi	di	dil		r13	r13d	r13w	r13b	
rbp	ebp	bp	bpl		r14	r14d	r14w	r14b	
rsp	esp	sp	spl		r15	r15d	r15w	r15b	
rip	eip	ip			rflags	eflags	flags		

- There are **16** 64-bit (`%rax ... %rsp`, `%r8 ... %r15`) general-purpose registers;
 - The **low-order 32, 16, and 8 bits** of each register **can be accessed** independently under other names.
- Almost any register can be used **to hold operands** for almost any logical and arithmetic operation, but **some have special or restricted uses**.
 - Register `%rsp` is reserved as the **stack pointer**.
 - Register `%rbp` is used as a frame pointer, i.e., **the base of the current stack frame**.
 - A few other instructions make implicit use of certain registers;
 - For example, the integer multiply and divide instructions require the `%rax` and `%rdx`.
 - `%rax` register is used for specifying the **syscall number** and it is also used **for storing the return value of functions**.
 - Calling functions: The **first six arguments of a function** are stored as follows: `%rdi`(arg0), `%rsi`(arg1), `%rdx`(arg2), `%rcx` (arg3), `%r8` (arg4), and `%r9` (arg5).
 - If there are more arguments than six they are stored on the stack.

- The instruction pointer register (`%rip`) **points to the next instruction to be executed**;
 - It **cannot be directly accessed by the programmer**, but is heavily used as the base for **position-independent code addressing**.
 - It is security feature.
 - It specifically supports **position-independent executables** (PIE), which are programs that work independently of where their code is loaded into memory
 - In a PIE, the operating system loads the program at varying locations: every time it runs, the program's functions and global variables have different addresses. This makes the program harder to attack (though not impossible)
 - Therefore, global variables are referenced relatively to the current value of the program counter (the `%rip` register in x86-64)

%rflags (I)

- When the ALU performs some operations, it **flags the results of these operations** in a special 64-bits register called it %rflags.



■ Carry Flag(CF):

- Set if the **arithmetic operation generates a carry or a borrow out** of the most significant bit of the result;
 - Cleared otherwise.
- The flag indicates **an overflow condition for unsigned-integer arithmetic**.

■ Zero Flag (ZF):

- Set if the **result is zero**;
 - Cleared otherwise.

■ Sign Flag (SF):

- Set **equal to the most significant bit of the result**, the sign bit for a signed integer.
 - 0 indicates a positive value;
 - 1 a negative value.

■ Overflow Flag (OF):

- Set if the integer **result is too large** excluding the sign-bit **to fit in the destination operand**.
 - Cleared otherwise.

■ The rules for turning on the carry flag in binary/integer math are two:

- 1 The carry flag is set if the addition of two numbers causes a carry out of the most significant (leftmost) bits added.

```
1111 + 0001 = 0000 (carry flag is turned on)
```

- 2 The carry (borrow) flag is also set if the subtraction of two numbers requires a borrow into the most significant (leftmost) bits subtracted.

```
0000 - 0001 = 1111 (carry flag is turned on)
```

■ Otherwise, the carry flag is turned off (zero).

```
0111 + 0001 = 1000 (carry flag is turned off [zero])
1000 - 0001 = 0111 (carry flag is turned off [zero])
```

- In unsigned arithmetic, watch the carry flag to detect errors.
- In signed arithmetic, the carry flag tells you nothing interesting.

- The rules for turning on the overflow flag in binary/integer math are two:

- 1 If the sum of two numbers with the sign bits off yields a result number with the sign bit on, the "overflow" flag is turned on.

0100 + 0100 = 1000 (overflow flag is turned on)

- 2 If the sum of two numbers with the sign bits on yields a result number with the sign bit off, the "overflow" flag is turned on.

1000 + 1000 = 0000 (overflow flag is turned on)

- Otherwise, the overflow flag is turned off (zero).

0100 + 0001 = 0101 (overflow flag is turned off)
0110 + 1001 = 1111 (overflow flag is turned off)
1000 + 0001 = 1001 (overflow flag is turned off)
1100 + 1100 = 1000 (overflow flag is turned off)

- In signed arithmetic, watch the overflow flag to detect errors.
- In unsigned arithmetic, the overflow flag tells you nothing interesting.
 - Mixed-sign addition never turns on the overflow flag.

- The x86-64 registers, memory and operations use the following data types (among others):

Data type	Suffix	Size (nr bytes)
byte	b	1
word	w	2
double (or long) word	l	4
quad word	q	8

- The **suffix** is used by the GNU assembler (gas) to specify appropriately-sized variants of instructions.

- Assembly language defines intrinsic data types, each of which describes a set of values that can be assigned to variables and expressions of the given type.
- The essential characteristic of each type is its **size in bits**.
 - `.octa` – 128 bits (16 bytes) integer
 - `.quad` – 64 bits (8 bytes) integer
 - `.long` – the same as `.int`
 - `.int` – 32 bits (4 bytes) integer
 - `.short` – 16 bits (2 bytes) integer
 - `.byte` – 8 bits (1 byte) integer
 - `.ascii` – string (with no automatic trailing zero byte)
 - `.asciz` – string automatically terminated by zero (The “z” stands for “zero”)
 - `.float` – floating point number (4 bytes)
 - `.double` – floating point number with double precision (8 bytes)
- The `.lcomm` and `.comm` directives allocate storage in the `.bss` section.

Statements

- Input to the assembler is a **text file**, with extension ".s", consisting of a **sequence of statements**.
- A statement consists of **tokens** separated by **whitespace** and terminated by either a newline character or a semicolon (;).
- **Whitespace** consists of spaces and tabs that are not contained in a string or comment.
- A statement can consist of a **comment**, started by #.
 - The comment is terminated by the newline that terminates the statement.
- An empty statement is one that contains nothing other than spaces, tabs, newlines or other similar characters.
 - Empty statements have no meaning to the assembler.

- A label **can be placed at the beginning of a statement.**
- During assembly, the label **is assigned the current value of the active location address** and serves as **an instruction operand.**
- A symbolic label consists of an identifier (or symbol) followed by a colon (:).
 - Symbolic labels must be defined only once.

■ Identifiers (symbols)

- An identifier is an arbitrarily-long **sequence of letters and digits**.
- The first character must be a letter; the underscore (`_`) and the period (`.`) are considered to be letters.
- Case is significant: uppercase and lowercase letters are different.

■ Keywords

- Keywords such as x86-64 instruction mnemonics (“opcodes”) and assembler directives are **reserved for the assembler and should not be used as identifiers**.

■ Numerical Constants

- Numbers can be integers or floating point ¹.
- Integers can be signed or unsigned, with signed integers represented in two's complement representation.

¹Out of scope

■ Numerical Constants (cont.)

■ Can be expressed

- Decimal integers begin with a non-zero digit followed by zero or more decimal digits (0–9).
- Binary integers begin with “0b” or “0B” followed by zero or more binary digits (0, 1).
- Octal integers begin with zero (0) followed by zero or more octal digits (0–7).
- Hexadecimal integers begin with “0x” or “0X” followed by one or more hexadecimal digits (0–9, A–F). Hexadecimal digits can be either uppercase or lowercase.

■ String Constants

- A string constant consists of a sequence of characters enclosed in double quotes (").

Directives

- Directives are commands that are part of the assembler syntax but are not related to the x64 processor instruction set.
- All assembler directives begin with a period (.).

- The general form of an assembly instruction is: **mnemonic operands**
 - Each mnemonic **opcode** represents a CPU instruction.
 - Instructions **can have a variable number of operands** and when more than one are separated by commas ", ".
 - For instructions with two operands, the first (lefthand) operand is the source operand, and the second (righthand) operand is the destination operand.
 - An operand specifies data being operated on or manipulated.
 - An operand has a type that can either be a register, a memory location, an immediate value or an address.
- The term addressing modes refers to the way in which the operand of an instruction is specified.

Addressing modes

- **Immediate mode:** the operand is specified by a \$ followed by an numerical constant value.
- **Register mode:** the operand is the name of a register.

- **Memory mode:** the operand is the value stored in a memory location, which is specified by an offset from the value in a register.
 - `offset (base, index, scale)` where `base` and `index` are registers, `scale` is a constant 1,2,4, or 8, and `offset` is a constant or symbolic label.
 - The effective address corresponding to this specification is $(base + index \times scale + offset)$
 - Any of the various fields may be omitted if not wanted; in effect, the omitted field contributes 0 to the effective address (except that `scale` defaults to 1).
 - **RIP-relative addressing:** this is new for x64 and allows accessing data tables and such in the code relative to the current instruction pointer, making position independent code easier to implement.
 - For example, we would write the address of a global value stored at location labeled `a` as `a(%rip)`, meaning that the assembler and linker should cooperate to compute the offset of `a` from the ultimate location of the current instruction.

- The instructions in the `mov` class **copy their source (S) values to their destinations (D)**.
 - The source operand designates a value that is immediate (I), stored in a register (R), or stored in memory (M).
 - The destination operand designates a location that is either a register or a memory address.
 - A move instruction cannot have both operands refer to memory locations.

Instruction	Description
<code>mov S, D</code>	Move source to destination
<code>movs S, D</code>	Move byte to word (sign extended)
<code>movz S</code>	Move byte to word (zero extended)

Unary Operations

Instruction	Description
inc <i>D</i>	Increment by 1
dec <i>D</i>	Decrement by 1
neg <i>D</i>	Arithmetic negation
not <i>D</i>	Bitwise complement

Binary Operations

Instruction	Description
add <i>S, D</i>	Add source to destination
sub <i>S, D</i>	Subtract source from destination
imul <i>S, D</i>	Multiply destination by source
xor <i>S, D</i>	Bitwise XOR destination by source
or <i>S, D</i>	Bitwise OR destination by source
and <i>S, D</i>	Bitwise AND destination by source

- Shift Operations
 - k is a numeric literal

Instruction	Description
sal/shl k, D	Left shift destination by k bits
sar k, D	Arithmetic right shift destination by k bits
shr k, D	Logical right shift destination by k bits

■ Special Arithmetic Operations

Instruction	Description
imulq <i>S</i>	Signed full multiply of %rax by <i>S</i> Result stored in %rdx:%rax
mulq <i>S</i>	Unsigned full multiply of %rax by <i>S</i> Result stored in %rdx:%rax
idivq <i>S</i>	Signed divide %rdx:%rax by <i>S</i> Quotient stored in %rax Remainder stored in %rdx
divq <i>S</i>	Unsigned divide %rdx:%rax by <i>S</i> Quotient stored in %rax Remainder stored in %rdx

Comparison and Test Instruction

Instruction	Description
cmp <i>S2</i> , <i>S1</i>	Set condition codes according to <i>S1</i> - <i>S2</i>
test <i>S2</i> , <i>S1</i>	Set condition codes according to <i>S1</i> & <i>S2</i>

Conditional Set Instructions

Instruction	Description	Condition
sete / setz <i>D</i>	Set if equal/zero	ZF
setne / setnz <i>D</i>	Set if not equal/nonzero	~ZF
sets <i>D</i>	Set if negative	SF
setns <i>D</i>	Set if nonnegative	~SF
setg / setnle <i>D</i>	Set if greater (signed)	~(SF^OF)&~ZF
setge / setnl <i>D</i>	Set if greater or equal (signed)	~(SF^OF)
setl / setnge <i>D</i>	Set if less (signed)	SF^OF
setle / setng <i>D</i>	Set if less or equal	(SF^OF) ZF
seta / setnbe <i>D</i>	Set if above (unsigned)	~CF&~ZF
setae / setnb <i>D</i>	Set if above or equal (unsigned)	~CF
setb / setnae <i>D</i>	Set if below (unsigned)	CF
setbe / setna <i>D</i>	Set if below or equal (unsigned)	CF ZF

■ Jump Instructions

Instruction	Description	Condition
<i>jmp Label</i>	Jump to label	
<i>je / jz Label</i>	Jump if equal/zero	ZF
<i>jne / jnz Label</i>	Jump if not equal/nonzero	~ZF
<i>js Label</i>	Jump if negative	SF
<i>jns Label</i>	Jump if nonnegative	~SF
<i>jg / jnle Label</i>	Jump if greater (signed)	~(SF^0F)&~ZF
<i>jge / jnl Label</i>	Jump if greater or equal (signed)	~(SF^0F)
<i>jl / jnge Label</i>	Jump if less (signed)	SF^0F
<i>jle / jng Label</i>	Jump if less or equal	(SF^0F) ZF
<i>ja / jnbe Label</i>	Jump if above (unsigned)	~CF&~ZF
<i>jae / jnb Label</i>	Jump if above or equal (unsigned)	~CF
<i>jb / jnae Label</i>	Jump if below (unsigned)	CF
<i>jbe / jna Label</i>	Jump if below or equal (unsigned)	CF ZF

Conditional Move Instructions

Instruction	Description	Condition
<i>cmove / cmovz S, D</i>	Move if equal/zero	ZF
<i>cmovne / cmovnz S, D</i>	Move if not equal/nonzero	~ZF
<i>cmovs S, D</i>	Move if negative	SF
<i>cmovns S, D</i>	Move if nonnegative	~SF
<i>cmovg / cmovnl S, D</i>	Move if greater (signed)	~(SF^0F)&~ZF
<i>cmovge / cmovnl S, D</i>	Move if greater or equal (signed)	~(SF^0F)
<i>cmovl / cmovnge S, D</i>	Move if less (signed)	SF^0F
<i>cmovle / cmovng S, D</i>	Move if less or equal	(SF^0F) ZF
<i>cmova / cmovnbe S, D</i>	Move if above (unsigned)	~CF&~ZF
<i>cmovae / cmovnb S, D</i>	Move if above or equal (unsigned)	~CF
<i>cmovb / cmovnae S, D</i>	Move if below (unsigned)	CF
<i>cmovbe / cmovna S, D</i>	Move if below or equal (unsigned)	CF ZF

Procedure Call Instruction

Instruction	Description
call <i>Label</i>	Push return address and jump to label
ret	Pop return address from stack and jump there

Stack

Instruction	Description
push <i>S</i>	Push source onto stack
pop <i>D</i>	Pop top of stack into destination

Convert

Instruction	Description
cwtl	Convert word in <code>%ax</code> to doubleword in <code>%eax</code> (sign-extended)
cltq	Convert doubleword in <code>%eax</code> to quadword in <code>%rax</code> (sign-extended)
cqto	Convert quadword in <code>%rax</code> to octoword in <code>%rdx:%rax</code> (sign-extended)

The 64 bit x86 C Calling Convention

Calling Convention

- Whenever you have assembler coded function invoked from C code **some issues arise**:
 - How are parameters passed to a function?
 - Can functions overwrite the values in a register, or does the caller expect the register contents to be preserved?
 - Where should local variables in a function be stored?
 - How should results be returned from functions?
- The C calling convention is based heavily on the use of the hardware-supported **stack**.
 - To understand the C calling convention, you should first make sure that you fully understand the `push`, `pop`, `call`, and `ret` instructions
- A **caller** is a function that calls another function;
- A **callee** is a function that was called.

63	31	15	8	7	0	
%rax	%eax	%ax	%ah	%al		Return value
%rbx	%ebx	%bx	%bh	%bl		Callee saved
%rcx	%ecx	%cx	%ch	%cl		4th argument
%rdx	%edx	%dx	%dh	%dl		3rd argument
%rsi	%esi	%si		%sil		2nd argument
%rdi	%edi	%di		%dil		1st argument
%rbp	%ebp	%bp		%bpl		Callee saved
%rsp	%esp	%sp		%spl		Stack pointer
%r8	%r8d	%r8w		%r8b		5th argument
%r9	%r9d	%r9w		%r9b		6th argument
%r10	%r10d	%r10w		%r10b		Caller saved
%r11	%r11d	%r11w		%r11b		Caller saved
%r12	%r12d	%r12w		%r12b		Callee saved
%r13	%r13d	%r13w		%r13b		Callee saved
%r14	%r14d	%r14w		%r14b		Callee saved
%r15	%r15d	%r15w		%r15b		Callee saved

The Caller's Rules (I)

- 1 Before calling a function, the caller should save the contents of certain registers that are designated **caller-saved**.
 - The caller-saved registers are `%r10`, `%r11`, and any registers that parameters are put into.
 - If you want the contents of these registers **to be preserved across the function call, push them onto the stack**.
- 2 To pass parameters to the function, we put up to six of them into registers (in order: `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8`, `%r9`).
 - If there are **more than six parameters** to the function, then **push the rest onto the stack in reverse order** (i.e. last parameter first)
 - Since the stack grows down, the first of the extra parameters (really the seventh parameter) parameter will be stored at the lowest address (this inversion of parameters was historically used to allow functions to be passed a variable number of parameters).

The Caller's Rules (II)

- 3 To call the function, use the `call` instruction.
 - This instruction places the **return address on top of the parameters on the stack**, and **branches to the subroutine code**.
- 4 After the function returns, (i.e. immediately following the call instruction) the caller must remove any additional parameters (beyond the six stored in registers) from stack.
 - This restores the stack to its state before the call was performed.
- 5 The caller can expect to find the **return value of the function** in the register `%rax`.
- 6 The caller restores the contents of caller-saved registers (`%r10`, `%r11`, and any in the parameter passing registers) by **popping them off of the stack**.
 - The caller can assume that no other registers were modified by the function.

- 1 Allocate **local variables by using registers or making space on the stack**.
 - Recall, the stack grows down, so to make space on the top of the stack, the stack pointer (`%rsp`) should be decremented.
 - The amount by which the stack pointer is decremented depends on the number of local variables needed
- 2 The values of any registers that are designated **callee-saved that will be used by the function must be saved**.
 - To save registers, **push them onto the stack**.
 - The callee-saved registers are `%rbx`, `%rbp`, and `%r12` through `%r15` (`%rsp` will also be preserved by the call convention, but need not be pushed on the stack during this step).
- 3 When the function is done, the **return value for the function** should be placed in `%rax` if it is not already there.

- 4 The function must **restore the old values of any callee-saved registers** (`%rbx`, `%rbp`, and `%r12` through `%r15`) that were modified.
 - The register contents are restored **by popping them from the stack**.
 - Note, the registers should be popped in the inverse order that they were pushed.
- 5 Next, we deallocate local variables.
 - The easiest way to do this is to add to `%rsp` the same amount that was subtracted from it in step 1.
- 6 Finally, we return to the caller by executing a `ret` instruction.
 - This instruction will find and remove the appropriate **return address from the stack**, pushed by `call` instruction