

ARQCP Course

Arquitetura de Computadores
Licenciatura em Engenharia Informática

2023/24
Paulo Baltarejo Sousa
pbs@isep.ipp.pt

ISEP INSTITUTO SUPERIOR
DE ENGENHARIA DO PORTO

Material and Slides

Some of the material/slides are adapted from various:

- Presentations found on the internet;
- Books;
- Web sites;
- ...

- 1 Data Transfer
- 2 Arithmetic Operations
- 3 Logic Operations
- 4 Control
- 5 `if, switch & loops`

Data Transfer

- `offset(base, index, scale)` where `base` and `index` are registers, `scale` is a constant 1,2,4, or 8, and `offset` is a constant or symbolic label.
 - The effective address corresponding to this specification is $(offset + R[base] + R[index] * scale)$.
 - Any of the various fields may be omitted if not wanted; in effect, the omitted field contributes 0 to the effective address (except that `scale` defaults to 1).

Syntax	Meaning
<code>0x104</code>	Address 0x104 (no \$)
<code>(%rax)</code>	What's in <code>%rax</code>
<code>4(%rax)</code>	What's in <code>%rax</code> , plus 4
<code>(%rax, %rdx)</code>	Sum of what's in <code>%rax</code> and <code>%rdx</code>
<code>4(%rax, %rdx)</code>	Sum of values in <code>%rax</code> and <code>%rdx</code> , plus 4
<code>(, %rcx, 4)</code>	What's in <code>%rcx</code> , times 4 (multiplier can be 1, 2, 4, 8)
<code>(%rax, %rcx, 2)</code>	What's in <code>%rax</code> , plus 2 times what's in <code>%rcx</code>
<code>8(%rax, %rcx, 2)</code>	What's in <code>%rax</code> , plus 2 times what's in <code>%rcx</code> , plus 8

- Usage: `mov[b|w|l|q] S, D`
 - Actually, `mov` instruction copies data
 - For most cases, the `mov` instructions will only update the specific register bytes or memory locations indicated by the destination operand
 - The only exception is that when `movl` has a register as the destination, it will also set the high-order 4 bytes of the register to 0
 - This exception arises from the convention, adopted in x86-64, that any instruction that generates a 32-bit value for a register also sets the high-order portion of the register to 0

```
# %rax = 0x0011223344556677
movb $-1,%al # %rax = 0x00112233445566FF
movw $-1,%ax # %rax = 0x001122334455FFFF
movl $-1,%eax # %rax = 0x00000000FFFFFFFF
movq $-1,%rax # %rax = 0xFFFFFFFFFFFFFFFF
```

- The regular `movq` instruction can **only have immediate source operands that can be represented as 32-bit two's-complement numbers**
 - This value is then sign extended to produce the 64-bit value for the destination
- The `movabsq` (move absolute quad word) instruction **can have an arbitrary 64-bit immediate value as its source operand** and can only have a register as a destination
 - Usage: `movabsq Imm, Reg`

```
movq $0x44556677,%rbx      # %rbx = 0x0000000044556677
movabsq $0x0011223344556677,%rax  # %rax = 0x0011223344556677
```

movq operand combinations

	Source	Dest	Src, Dest	C Analog
movq	Imm	Reg	movq \$0x4,%rax	temp = 0x4;
		Mem	movq \$-147,(%rax)	*p = -147;
	Reg	Reg	movq %rax,%rdx	temp2 = temp1;
		Mem	movq %rax,(%rdx)	*p = temp;
	Mem	Reg	movq (%rax),%rdx	temp = *p;

■ Cannot do **memory-memory transfer with a single instruction.**

■ The **first six arguments of a function** are stored as follows:

Operand size (bits)	Arguments					
	1st	2nd	3rd	4th	5th	6th
64	%rdi	%rsi	%rdx	%rcx	%r8	%r9
32	%edi	%esi	%edx	%ecx	%r8d	%r9d
16	%di	%si	%dx	%cx	%r8w	%r9w
8	%dil	%sil	%dl	%cl	%r8b	%r9b

■ If there are **more than six** they are stored on the **stack**.

void swap(long* px, long *py) (I)

```
void swap(long* px, long* py)
{
    long t0 = *px;
    long t1 = *py;
    *px = t1;
    *py = t0;
}
```

```
#include <stdio.h>

int main(void) {
    long x = 20;
    long y = 50;
    printf("before swap: x= %ld: y= %ld\n", x,y);
    swap(&x, &y);
    printf("after swap: x= %ld: y= %ld\n", x,y);
    return 0;
}
```

```
.section .text
.global swap
swap:
    # prologue

    movq (%rdi), %rax
    movq (%rsi), %rdx
    movq %rdx, (%rdi)
    movq %rax, (%rsi)

    # epilogue
    ret
```


■ The `movs` class fills out the remaining bytes of the **destination with bit sign**

■ Usage: `movs [bw|bl|bq|wl|wq|lq] S, D`

■ The `movz` class fills out the remaining bytes of the **destination with zeros**

■ Usage: `movz [bw|bl|bq|wl|wq] S, D`

- Note the absence of an explicit instruction to zero-extend a 4-byte source value to an 8-byte destination [lq]
- This type of data movement can be implemented using a `movl` instruction having a register as the destination



Arithmetic Operations

void sumstore(long x, long y, long *dest) (I)

```
void sumstore(long x, long y, long *dest)
{
    long t = x + y;
    *dest = t;
}
```

```
#include <stdio.h>

int main(void) {
    long x = 0;
    sumstore(20, 30, &x);
    printf("x= %ld\n", x);
    return 0;
}
```

```
.section .text
.global sumstore
sumstore:
    # prologue

    movq %rdi, %rbx
    addq %rsi, %rbx
    movq %rbx, (%rdx)

    # epilogue
    ret
```

- Most of the operations are given as instruction classes, as they can have different variants with different operand sizes, specified by suffixes (b, w, l, and q)
- However, **in an operation the operands must be of the same size**, according **to the instruction suffix**.
 - If not, the **type conversion has higher precedence than arithmetic operations**
- Moving **from a smaller to a larger** data size can involve either **sign extension** (for signed values) or **zero extension** (for unsigned values)
- There are two `mov` instructions that can be used to copy a smaller source to a larger destination:
 - `movz` fills the remaining bytes with zeros
 - `movs` fills the remaining bytes by sign-extending the most significant bit in the source.
 - The source could be from memory or a register, and the destination is a register.

void sumstore2(char x, short y, long *dest) (I)

```
void sumstore2(char x, short y, long *dest)
{
    long t = x + y;
    *dest = t;
}
```

```
#include <stdio.h>

int main(void) {
    long x = 0;
    sumstore2(20, 30, &x);
    printf("x= %ld\n", x);
    return 0;
}
```

```
.section .text
.global sumstore2
sumstore2:
    # prologue

    movsbq %dil, %rbx
    movswq %si, %rcx
    addq %rbx, %rcx
    movq %rcx, (%rdx)

    # epilogue
    ret
```

Large Multiplication

- Multiplying 64-bit numbers can produce a 128-bit result.
- How does x86-64 support this with only 64-bit registers?
 - **Join two 64-bit registers to hold a 128-bit result.**

Instruction	Effect	Description
imulq S	$R[\%rdx]:R[\%rax] \leftarrow S * R[\%rax]$	Signed full multiply
mulq S	$R[\%rdx]:R[\%rax] \leftarrow S * R[\%rax]$	Unsigned full multiply

- The high-order 64 bits of the **dividend** are in `%rdx`, and the **low-order 64 bits** are in `%rax`.
- The **divisor** is the operand to the instruction.
- The **quotient** is stored in `%rax`, and the **remainder** in `%rdx`.

Instruction	Effect	Description
<code>idivq S</code>	$R[\%rax] \leftarrow R[\%rdx]:R[\%rax] / S$ $R[\%rdx] \leftarrow R[\%rdx]:R[\%rax] \bmod S$	Signed divide
<code>divq S</code>	$R[\%rax] \leftarrow R[\%rdx]:R[\%rax] / S$ $R[\%rdx] \leftarrow R[\%rdx]:R[\%rax] \bmod S$	Unsigned divide
<code>cqto</code>	$R[\%rdx]:R[\%rax] \leftarrow \text{SignExtend}(R[\%rax])$	Convert to oct word

- Most **division** uses only 64 bit dividends.
- The `cqto` instruction sign-extends the 64-bit value in `%rax` into `%rdx` to fill both registers with the dividend, as the division instruction expects.

```
void divide(long dividend, long divisor, long* quotient, long* remainder) (I)
```

```
void divide(long dividend, long divisor, long* quotient, long* remainder)
{
    *quotient = dividend / divisor;
    *remainder = dividend % divisor;
}
```

```
#include <stdio.h>

int main(void) {
    long q = 0, r = 0;
    divide(100, 30, &q, &r);
    printf("q= %ld, r= %ld\n", q, r);
    return 0;
}
```

```
void divide(long dividend, long divisor, long*
quotient, long* remainder)(ll)
```

```
.section .text
.global divide
divide:
# prologue

    movq %rdi, %rax
    movq %rdx, %rdi
    cqto
    idivq %rsi
    movq %rax, (%rdi)
    movq %rdx, (%rcx)


# epilogue
    ret
```

Logic Operations

- Logical instructions **play a pivotal role in manipulating data at the bit level.**
- These instructions **allow for precise control over individual bits**, making them essential for various tasks.
- One of the most common uses of logical operations is **bit masking**.
- Bit masking **is a technique in programming used to test or modify the states of the bits** of a given data.

Bit masking (II)

- The main usage of bitwise logical instructions is: **to set, to clear, to invert, to isolate** some selected bits in the Destination operand.
- To do this, a Source bit pattern known as **a mask** is constructed.
- The mask bits **are chosen so that the selected bits are modified in the desired manner when an instruction is executed.**



- The Mask bits **must be constructed based on the properties of and, or, and xor operations.**

- The `and` instruction:

- Can be used to clear specific Destination bits while preserving the others.
- A zero mask bit clears the corresponding Destination bit;
- A one mask bit preserves the corresponding Destination bit.

- The `or` instruction:

- Can be used to set specific Destination bits while preserving the others.
- A one mask bit sets the corresponding Destination bit;
- A zero mask bit preserves the corresponding Destination bit.

- The `xor` instruction:


- Can be used to invert specific Destination bits while preserving the others.
- A one mask bit inverts the corresponding Destination bit;
- A zero mask bit preserves the corresponding Destination bit.

■ Swap values

```
xorq %rdi, %rsi
xorq %rsi, %rdi
xorq %rdi, %rsi
```

■ Clear a register

```
andq $0, %rax
```



- Arithmetic Shift operations **can be used for dividing or multiplying** an integer variable **by powers of 2**.
 - **Multiplication by left shift:** the result of a Left Shift operation is a multiplication by 2^n , where n is the number of shifted bit positions.
 - **Division by right shift:** The result of a Right Shift operation is a division by 2^n , where n is the number of shifted bit positions.

Control

- The CPU contains a 64-bit register, `%rflags`, **where individual bits (flags) are set or cleared as a consequence of arithmetic and other operations.**
- `%rflags` register holds information **about the result of the most recent instructions that has affected flags.**
- These can be **updated and read to influence what to do next.**
- Most common flags, **condition codes (cc)**, are:
 - **Carry flag (CF).**
 - The most recent operation generated a carry out of the most significant bit.
 - Used to detect overflow for unsigned operations.
 - **Zero flag (ZF).**
 - The most recent operation yielded zero.
 - **Sign flag (SF).**
 - The most recent operation yielded a negative value.
 - **Overflow flag (OF).**
 - The most recent operation caused a two's-complement overflow-either negative or positive.
- **They can be set implicitly or explicitly**

Condition codes: Implicit setting

- Implicitly set (**think of it as side effect**) by arithmetic operations
- `add a, b`
 - `t = a + b` setting destination to `t`
 - Sets condition codes based on value of `a + b`
 - CF set if carry out from most significant bit (unsigned overflow)
 - ZF set if `t == 0`
 - SF set if `t < 0` (as signed)
 - OF set if two's-complement (signed) overflow, whenever
`(a>0 && b>0 && t<0) || (a<0 && b<0 && t>=0)`

Condition codes: Explicit setting

■ `cmp a, b`

- The `cmp` instruction is like the subtraction instruction, but it does not store the result anywhere.
- Sets condition codes based on value of $b - a$
 - CF set if carry out from most significant bit (unsigned overflow)
 - ZF set if $a == b$
 - SF set if $b - a < 0$ (as signed)
 - OF set if two's-complement (signed) overflow, whenever
 $(a < 0 \ \&\& \ b > 0 \ \&\& \ (b - a) < 0) \ || \ (a > 0 \ \&\& \ b < 0 \ \&\& \ (b - a) > 0)$

■ `test a, b`

- The `test` instruction is like the `and` instruction, but it does not store the result anywhere.
- Sets condition codes based on value of $a \ \& \ b$
 - ZF set if $a \ \& \ b == 0$
 - SF set if $a \ \& \ b < 0$
 - OF and CF are always set to 0
- Typically, the same operand is repeated (e.g., `testl %eax, %eax` to see whether `%eax` is negative, zero, or positive), or one of the operands is a mask indicating which bits should be tested.

- `cmp a, b`
 - Sets condition codes based on value of $b - a$

Suffix <i>cc</i>	Condition tested	Meaning after <code>cmp</code>
e	ZF	equal to zero
ne	~ZF	not equal to zero
s	SF	negative
ns	~SF	non-negative
g	~(SF^OF) & ~ZF	greater (signed >)
ge	~(SF^OF)	greater or equal (signed >=)
l	SF^OF	less (signed <)
le	(SF^OF) ZF	less or equal (signed <=)
a	~CF & ~ZF	above (unsigned >)
ae	~CF	above or equal (unsigned >=)
b	CF	below (unsigned <)
be	CF ZF	below or equal (unsigned <=)

Suffix <i>cc</i>	Meaning	<code>cmp a, b</code>	<code>test a, b</code>
e	Equal	<code>b == a</code>	<code>b & a == 0</code>
ne	Not equal	<code>b != a</code>	<code>b & a != 0</code>
s	Sign (negative)	<code>b - a < 0</code>	<code>b & a < 0</code>
ns	(non-negative)	<code>b - a >= 0</code>	<code>b & a >= 0</code>
g	Greater	<code>b > a</code>	<code>b & a > 0</code>
ge	Greater or equal	<code>b >= a</code>	<code>b & a >= 0</code>
l	Less	<code>b < a</code>	<code>b & a < 0</code>
le	Less or equal	<code>b <= a</code>	<code>b & a <= 0</code>
a	Above (unsigned >)	<code>b > a</code>	<code>b & a > 0U</code>
b	Below (unsigned <)	<code>b < a</code>	<code>b & a < 0U</code>

- `cmp a, b` is like `sub (b-a)`, `test` is like `and (a & b)`
- Result is not stored anywhere, they only set condition code bits

- There are **three common instruction types that use condition codes (cc)**:
 - `setcc` instructions **conditionally set a byte to 0 or 1**
 - `cmovcc` instructions **conditionally move data**
 - `jcc` instructions **conditionally jump to a different next instruction**

- Each `setcc` instruction **has a designated destination:**
 - **Byte length register**
 - Do not alter remaining bytes in register
 - **Single-byte memory location**

```
int less(int a, int b){  
    return ( a < b );  
}
```

```
less:  
    cmpl %esi, %edi  
    setl %al  
    movzbl %al, %eax  
    ret
```

- The `cmovcc` instructions **check the condition codes and perform or not a move operation**.
- If the **condition is not satisfied**, a **move is not performed** and **execution continues** with the instruction following the `cmovcc` instruction.

```
int equals(int a, int b){  
    int x = 0;  
    if ( a == b ) {  
        x = 1;  
    }  
    return x;  
}
```

```
equals:  
    movl $0, %eax  
    cmpl %esi, %edi  
    cmovl $1, %eax  
    ret
```

■ The `jmp` instruction **jumps to another instruction** in the assembly code (“Unconditional Jump”).

■ `jmp Label` (Direct Jump)

```
jmp end
...
end:...
```

■ `jmp *Operand` (Indirect Jump)

```
jmp %rax # jump to instruction at address in %rax
```

Conditional `jcc`

- The `jcc` instruction **jumps conditionally (if certain conditions are true) to another instruction** in the assembly code.
- The `test` and `cmp` instructions are combined with the conditional and unconditional jump instructions **to implement most relational and logical expressions and all control structures**.

if, switch & loops

- The general form of an `if...else` statement in C is:

```
if (test-expr)
    then-statement
else
    else-statement
```

- where `test-expr` is **an integer expression that evaluates either to 0** (interpreted as meaning “false”) **or to a nonzero value** (interpreted as meaning “true”).
 - **Only one of the two branch statements** (`then-statement` or `else-statement`) **is executed**.
- `if...else` statement in assembly implementation typically adheres to:

```
t = test-expr;
if (!t)
    goto false;
then-statement
goto done;
false:
    else-statement
done:
```

if...else (II)

```
int equals (int a, int b){
    int x;
    if ( a == b ) {
        x = 1;
    }else{
        x = 0;
    }
    return x;
}
```

```
equals:
    cmpl %esi, %edi
    jne false
    movl $1, %eax
    jmp done
false:
    movl $0, %eax
done:
    ret
```

```
int equals(int a, int b){
    int x = 0;
    if ( a == b ) {
        x = 1;
    }
    return x;
}
```

```
equals:
    movl $0, %eax
    cmpl %esi, %edi
    jne false
    movl $1, %eax
false:
    ret
```

if...else (III)

```
int check(int x, int y){
    if (x < 3 && x == y) {
        return 1;
    } else {
        return 2;
    }
}
```

```
check:
    cmpl $3, %edi
    setl %dl
    cmpl %esi, %edi
    sete %al
    testb %al, %dl
    je false
    movl $1, %eax
    jmp done
false:
    movl $2, %eax
done:
    ret
```

- A `switch` statement provides a multi-way branching capability based on the value of an integer index.
- They are particularly useful when dealing with tests where there can be a large number of possible outcomes.
- An efficient implementation using a data structure called a **jump table**.
 - A jump table is an array where entry `i` is the address of a code segment implementing the action the program should take when the switch index equals `i`.
 - The code performs an array reference into the jump table using the switch index to determine the target for a jump instruction.
 - The advantage of using a jump table over a long sequence of if-else statements is that the time taken to perform the switch is independent of the number of switch cases.

switch (II)

Switch Form

```
switch (x) {  
  case val_0:  
    Block 0  
  case val_1:  
    Block 1  
    . . .  
  case val_n-1:  
    Block n-1  
}
```

Jump Table

JTab:

Targ0
Targ1
Targ2
•
•
•
Targn-1

Jump Targets

Targ0: Code Block 0

Targ1: Code Block 1

Targ2: Code Block 2

•

•

•

Targn-1: Code Block n-1

Approximate Translation

```
target = JTab[x];  
goto target;
```


switch (III)

C code:

```
switch (x) {  
  case 1: <code> break;  
  case 2: <code>  
  case 3: <code> break;  
  case 5:  
  case 6: <code> break;  
  case 7: <code> break;  
  default: <code>  
}
```

Use the jump table when $x \leq 7$:

```
if (x <= 7)  
  target = JTab[x];  
  goto target;  
else  
  goto default;
```



```
int switcher(int a, int b, int c)
{
    int answer;
    switch(a) {
        case 5:
            c = b ^ 15;
            /* Fall through */
        case 0:
            answer = c + 112;
            break;
        case 2:
        case 7:
            answer = (c + b) << 2;
            break;
        case 4:
            answer = a;
            break;
        default:
            answer = b;
    }
    return answer;
}
```

```
.L7:
.long .L3 #0
.long .L2 #1
.long .L4 #2
.long .L2 #3
.long .L5 #4
.long .L6 #5
.long .L2 #6
.long .L4 #7
```

```
switcher:
    movl %edi, %eax
    cmpl $7, %eax
    ja .L2
    jmp *.L7(, %eax, 4)
.L2:
    movl %esi, %eax
    jmp .L8
.L5:
    movl $4, %eax
    jmp .L8
.L6:
    movl %esi, %eax
    xorl $15, %eax
    movl %eax, %edx
.L3:
    movl %edx, %eax
    addl $112, %eax
    jmp .L8
.L4:
    movl %edx, %eax
    addl %esi, %eax
    sall $2, %eax
.L8:
```

Loops

- C provides several looping constructs—namely, `do-while`, `while`, and `for`.
- Combinations of **conditional tests and jumps are used to implement the effect of loops**.
- Most compilers **generate loop code based on the `do-while` form of a loop**, even though this form is relatively uncommon in actual programs.
- Other loops are transformed into `do-while` form and then compiled into machine code.

- The general **form of a do-while statement** is as follows:

```
do
    body-statement
while (test-expr);
```

- The effect of the loop is to repeatedly execute `body-statement`, evaluate `test-expr`, and continue the loop if the evaluation result is nonzero.
- The `body-statement` is executed at least once.
- This general form can be translated into conditionals and goto statements as follows:

```
loop:
    body-statement
    t = test-expr;
    if (t)
        goto loop;
```

do-while (II)

```
void do_while_loop(int n) {  
    int i = 0;  
    do{  
        i++;  
    } while (i < n);  
}
```

```
do_while_loop:  
    movl $0, %eax  
loop:  
    addl $1, %eax  
    cmpl %edi, %eax  
    jl loop  
    ret
```

- The general **form of a while statement** is as follows:

```
while (test-expr)
  body-statement
```

- It differs from `do-while` in that `test-expr` is evaluated and the loop is potentially terminated before the first execution of `body-statement`.
- There are a number of ways to translate a while loop into machine code.
- One common approach is to transform the code into a `do-while` loop by using a conditional branch to skip the first execution of the body if needed:

```
if (!test-expr)
  goto done;
do
  body-statement
while (test-expr);
done:
```

```
t = test-expr;
if (!t)
  goto done;
loop:
  body-statement
  t = test-expr;
  if (t)
    goto loop;
done:
```

```
goto test;
loop:
  body-statement
test:
  t = test-expr;
  if (t)
    goto loop;
```

while (II)

```
void while_loop(int n) {  
    int i = 0;  
    while (i < n) {  
        i++;  
    }  
}
```

```
while_loop:  
    movl $0, %eax  
    jmp test  
loop:  
    addl $1, %eax  
test:  
    cmpl %edi, %eax  
    jl loop  
ret
```

- The general **form of a for statement** is as follows:

```
for (init-expr; test-expr; update-expr)
    body-statement
```

- The C language standard states that the behavior of such a loop is identical to the following code, which uses a `while` loop:

```
init-expr;
while (test-expr) {
    body-statement
    update-expr;
}
```

```
init-expr;
if (!test-expr)
    goto done;
do {
    body-statement
    update-expr;
} while (test-expr);
done:
```

```
init-expr;
goto test;
loop:
    body-statement
    update-expr;
test:
    t = test-expr;
    if (t)
        goto loop;
```

for (II)

```
void for_loop(int n) {  
    for (int i = 0; i < n; i++) {  
        // body  
    }  
}
```

```
for_loop:  
    movl $0, %eax  
    jmp test  
loop:  
    addl $1, %eax  
test:  
    cmpl %edi, %eax  
    jl loop  
    ret
```

loop instruction

- Loop according to counter register (`%rcx`).

```
loop label
```

- Where, `label` is the target label that identifies the target instruction as in the jump instructions.
- The `loop` instruction assumes that the **counter register contains the loop count**.
- When the **loop instruction is executed**, the **counter register is decremented and the control jumps to the target label**, until the **counter register value reaches zero**.

```
movq $10, %rcx
iter:
loop iter
```

Conditional Loop Instructions

- `loope destination and loopz destination`
- Logic:
 - $\%rcx \leftarrow \%rcx - 1$
 - If $\%rcx > 0$ and $ZF=1$, jump to destination
- `loopne destination and loopnz destination`
- Logic:
 - $\%rcx \leftarrow \%rcx - 1$
 - If $\%rcx > 0$ and $ZF=0$, jump to destination