

ARQCP Course

Arquitetura de Computadores
Licenciatura em Engenharia Informática

2023/24
Paulo Baltarejo Sousa
`pbs@isep.ipp.pt`

ISEP INSTITUTO SUPERIOR
DE ENGENHARIA DO PORTO

Material and Slides

Some of the material/slides are adapted from various:

- Presentations found on the internet;
- Books;
- Web sites;
- ...

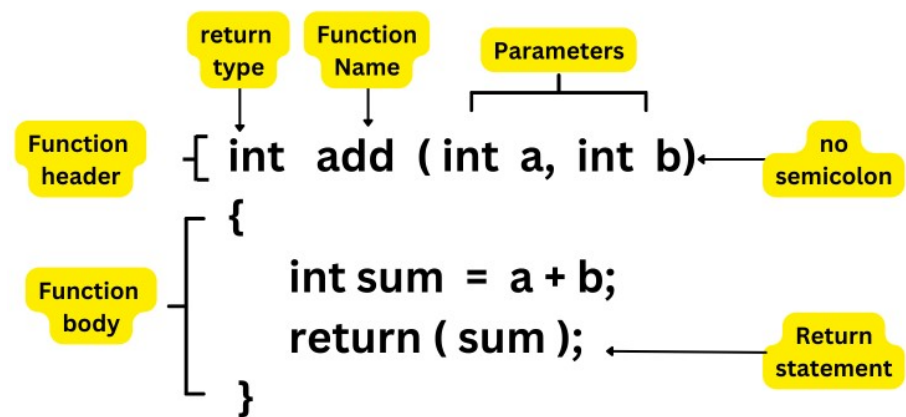
- 1 Functions
- 2 Stack
- 3 Calling conventions
- 4 Local storage
- 5 Stack Frames
- 6 Memory Errors Exploitation

Functions

What is?


- In computer programming, a **function** is a block of program instructions that performs a specific task, packaged as a **unit** (and identified by a **name**).
- Functions may be defined within **programs**, or separately in **libraries** that can be used by many programs.
 - A function may be called a **routine**, **subprogram**, **subroutine**, or **procedure**;
 - In object-oriented programming, it may be called a **method**.

Function: Definition



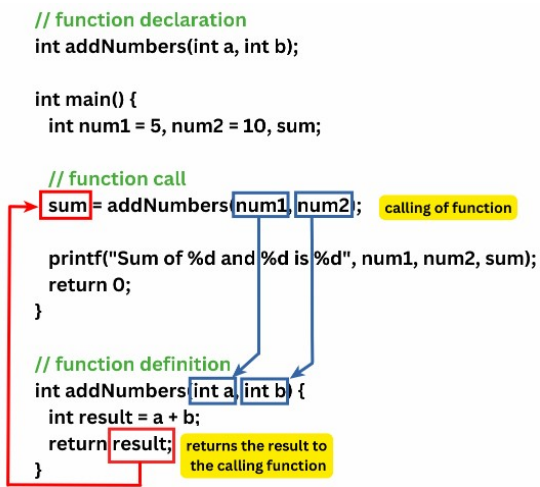
Function: Invocation

- A function is coded so that it can be **started/called several times and from several places** during one execution of the program and then **branch back (return) to the next instruction after the call**, once the function's **task is done**.
- A function's operation occurs only **when it is called**.



Function: Input & Output

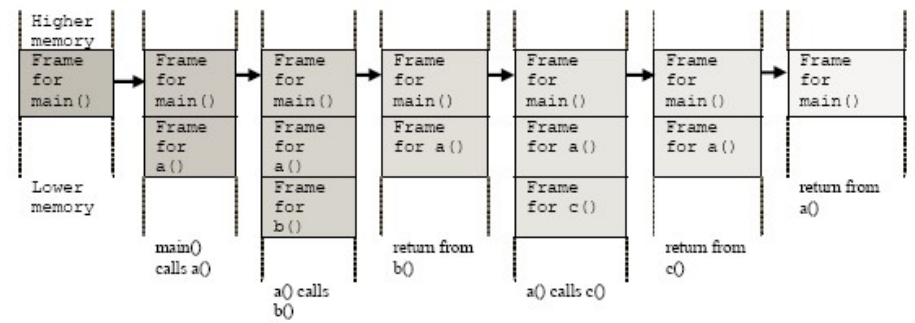
- Functions usually **take in** (parameters) data, **process** it, and **return** a result.



Function: Execution

```
int b() {  
    return 0;  
}  
  
int main() {  
    int x = a();  
    return 0;  
}
```

```
int c() {  
    return 0;  
}  
  
int a() {  
    int x = b();  
    int y = c();  
    return x+y;  
}
```



■ Passing control

- Invoking (calling) a function
- Return to the next instruction after the call

■ Passing data

- Function arguments
- Return value

■ Memory management

- Allocate during function execution
- Deallocate upon return

- These mechanisms are implemented with **register** and **stack** support


Application Binary Interface (ABI)

- An Application Binary Interface (ABI) is a **set of rules and conventions** that dictate how binary code or machine code communicates and interacts with other binary code, particularly in the context of software libraries, operating systems, or hardware.
- The ABI defines **data structures, calling conventions, register usage**, and other low-level details that ensure compatibility and interoperability between different software components.
- It acts as an **interface between high-level programming languages and the machine code**, allowing programs written in various languages to work together seamlessly.
- ABIs are crucial for binary compatibility and the proper functioning of software systems.

Stack

Virtual Memory Layout

- Memory viewed as **array of bytes** and different regions have different purposes




Stack it is used to handle functions.

As a program runs, calling one function after another, it continuously pushes data onto the stack and pops data off the stack, according to last in, first out (LIFO) heuristic.

For each function call it creates a stack frame.

`%rbp` and `%rsp` registers hold bottom and top addresses of the current stack frame, respectively.

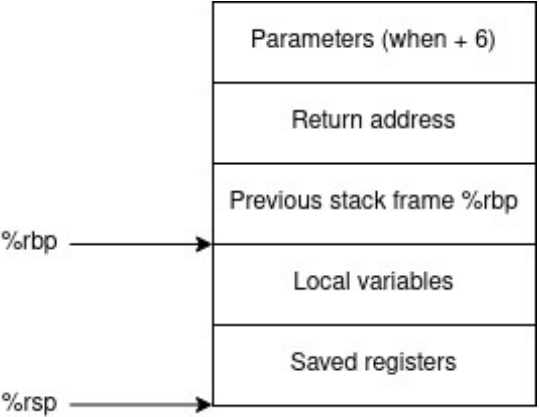


- **Whenever a function finishes** (the last one, which stack frame is at the top of the stack) its execution, **the stack frame is destroyed** and **%rbp and %rsp are updated**.



Stack Frame

- Stack frames are data structures that store information about function calls, such as parameters, local variables, return addresses, and saved registers.



Instruction: call

- call label
 - Push **Return address** on stack.
 - The **Return address** is the address of the next instruction right after call
 - %rsp is decremented by 8.
 - Sets the %rip register with the address represented by label

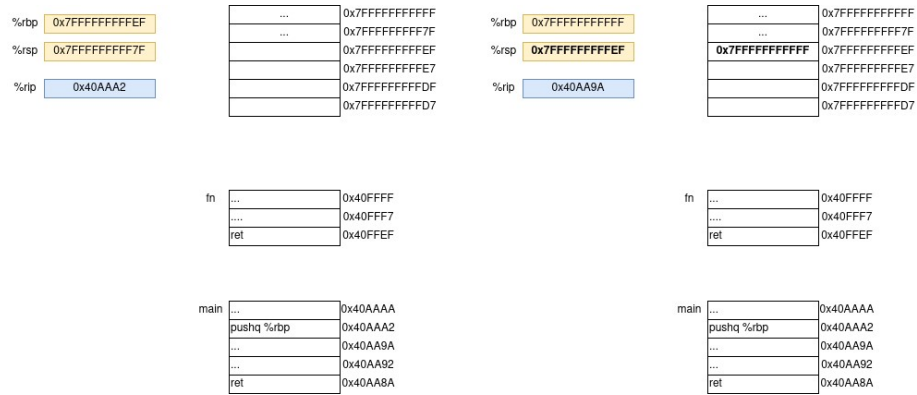


- ret instruction **pops the value off the stack and set %rip register with popped value.**
 - It should be the **Return address** previously pushed by call instruction.
 - %rsp is incremented by 8.



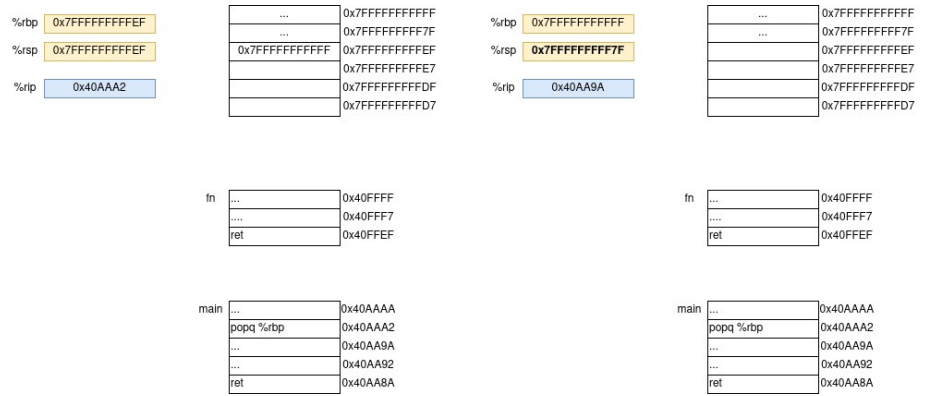
Instruction: pushq

- pushq *S*
 - Pushes *S* onto stack.
 - `%rsp` is decremented by 8.



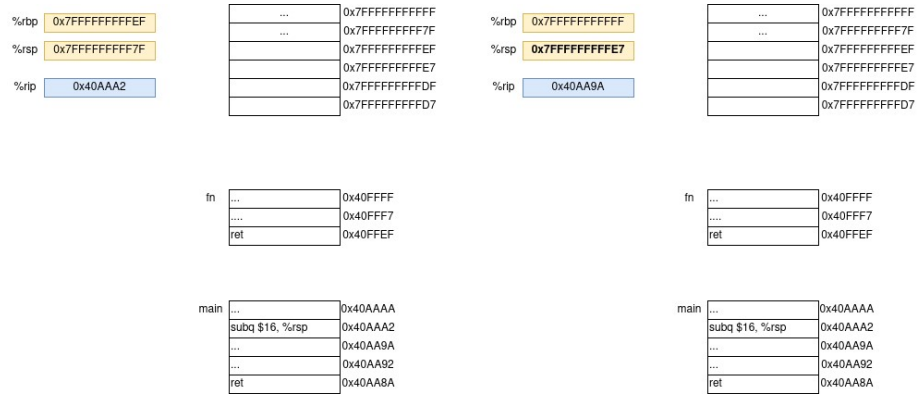
Instruction: popq

- popq *D*
 - Pop top of stack into *D*
 - %rsp* is incremented by 8.



Instruction: Increasing Stack

■ `subq $S, %rsp`



fn

...

...

ret

0x40FFFF

0x40FFF7

0x40FFEF

fn

...

....

ret

0x40FFFF

0x40FFF7

0x40FFEF

main

...

subq \$16, %rsp

...

...

ret

0x40AAAA

0x40AAA2

0x40AA9A

0x40AA92

0x40AA8A

main

...

subq \$16, %rsp

...

...

ret

0x40AAAA

0x40AAA2

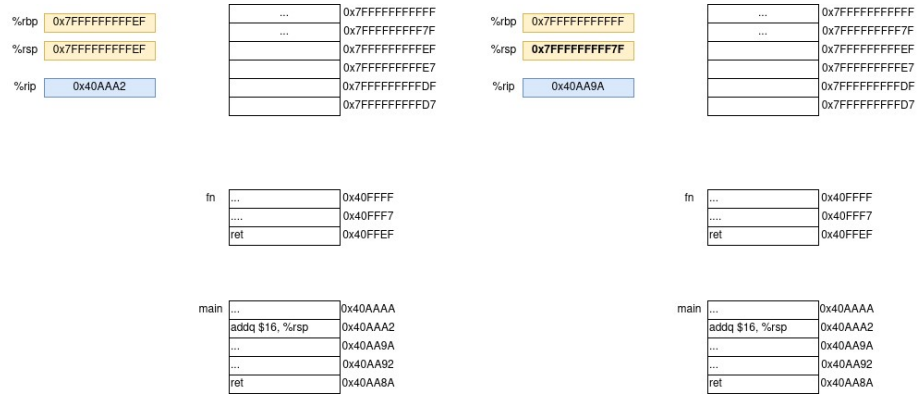
0x40AA9A

0x40AA92

0x40AA8A

Instruction: Decreasing Stack

addq \$, %rsp



Calling conventions

Calling Convention

- Calling conventions describe the **conventions or norms that functions use when it calls another function and when a function returns to its caller function.**
- The calling convention is based heavily on the use of the **stack** and **registers**.
- A **caller** is a function that calls another function;
- A **callee** is a function that was called.

- To **pass parameters to function**, we put **up to six of them into registers** (in order: %rdi, %rsi, %rdx, %rcx, %r8, %r9).
 - If there are **more than six parameters** to the function, then **push the rest onto the stack in reverse order** (i.e. last parameter first)

```
void f2( long x1,long x2,long x3,long x4,
        long x5,long x6,long x7,long x8);
```

%rax		%r8	x5	...	0x7FFFFFFFFFFF
%rbx		%r9	x6	...	0x7FFFFFFFFF7F
%rcx	x4	%r10		x8	0x7FFFFFFFFFEF
%rdx	x3	%r11		x7	0x7FFFFFFFFFE7
%rdi	x1	%r12			0x7FFFFFFFFDFD
%rsi	x2	%r13			0x7FFFFFFFFFD7
%rbp	0x7FFFFFFFFFFF	%r14			
%rsp	0x7FFFFFFFFFE7	%r15			

```
f1:
...
movq x1(%rip), %rdi
movq x2(%rip), %rsi
movq x3(%rip), %rdx
movq x4(%rip), %rcx
movq x5(%rip), %r8
movq x6(%rip), %r9
movq x8(%rip), %r10
pushq %r10
movq x7(%rip), %r10
pushq %r10
call f2
...
ret
```

Returning data to caller

- To return a value, the callee **stores it into %rax**
 - This is the reason for **a function returning only one value**

```
int x = 10;
long f2(){
    return 1;
}
void f1(){
    x += f2();
    ...
}
```

```
x: .int 10

f2:
...
    movq $1, %rax
...
    ret
f1:
...
    movq x(%rip), %r10
    call f2
    addq %rax, %r10
    ret
```

```
f1:
...
movq $50, %rbx
movq $200, %rdx
call f2
addq %rbx, %rdx
...
ret
```

```
f2:
...
movq $10, %rdx
...
addq $20, %rdx
...
ret
```

- Contents of register `%rdx` overwritten by `f2`
- This is an issue, that requires some **coordination mechanism**

```
f1:
...
movq $50, %rbx
movq $200, %rdx
pushq %rdx
call f2
popq %rdx
addq %rbx, %rdx
...
ret
```

```
f2:
...
movq $10, %rdx
...
addq $20, %rdx
...
ret
```

- **Caller saves %rdx before call and restores it after return**
 - Caller-saved registers **can be modified by any function**
 - Since the **called function is free to alter these registers**, it is incumbent upon the **caller to first save the data before it makes the call**

■ Caller save

- Caller saves (on Stack) temporary values before the call
- Caller restores (from Stack) them after returning from the callee

■ Callee save

- Callee saves (on Stack) temporary values before using
- Callee restores (from Stack) them before returning to caller

■ Which registers are **caller-save** or **callee-save**?

Register Saving Convention (II)

%rax	Return value - Caller saved
%rbx	Callee saved
%rcx	Argument #4 - Caller saved
%rdx	Argument #3 - Caller saved
%rsi	Argument #2 - Caller saved
%rdi	Argument #1 - Caller saved
%rsp	Stack pointer
%rbp	Callee saved

%r8	Argument #5 - Caller saved
%r9	Argument #6 - Caller saved
%r10	Caller saved
%r11	Caller Saved
%r12	Callee saved
%r13	Callee saved
%r14	Callee saved
%r15	Callee saved

■ These registers can be modified by **callee** function.

```
f1:
  pushq %rax
  pushq %rdi
  pushq %rsi
  pushq %rdx
  pushq %rcx
  pushq %r8
  pushq %r9
  pushq %r10
  pushq %r11
  call f2
  popq %r11
  popq %r10
  popq %r9
  popq %r8
  popq %rcx
  popq %rdx
  popq %rsi
  popq %rdi
  popq %rax
  ...
  ret
```

Return value

Arguments

Caller-saved
temporaries

%rax

%rdi

%rsi

%rdx

%rcx

%r8


%r9

%r10

%r11

Callee-saved registers

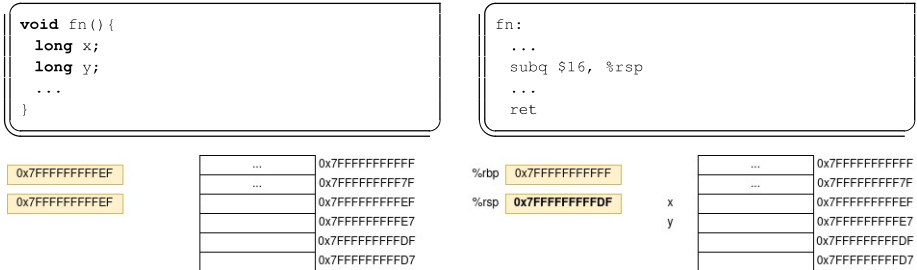
- **Callee must save and restore** (could be being used by caller)



- `%rbp`
 - May be used as **stack frame pointer**
- `%rsp`
 - Special form of **callee save**
 - **Restored to original value** upon exit from function.

Local storage

- **Local variables are allocated onto Stack**
- It simply **subtracting the number of bytes required by each variable from the `%rsp`**.
- This **does not store any data in the variables**, it simply sets aside memory that we can use.
- There are **no labels in this area of memory**




- Accessing to **local variable** could be done using `%rsp` as **anchor**.
- Recall, there are no labels in this area of memory

```
void fn(){
    long x;
    long y;
    ...
    x = 10;
    y = 20;
}
```

```
fn:
    ...
    subq $16, %rsp
    ...
    movq $20, (%rsp)
    movq $10, 8(%rsp)

    ret
```



Stack Frames

- Stack frames **only exist at run-time**.
- They are used **to handle the function calls**.
- Contents:
 - Local variables
 - Return information
 - Temporary space
- Management
 - Space allocated when enter function
 - “Set-up” code (prologue)
 - Deallocated when return
 - “Finish” code (epilogue)

Stack frames: Prologue and Epilogue (I)

```
function:
#Prologue
pushq %rbp    # Save old %rbp
movq %rsp,%rbp # Set %rbp as frame pointer
...
#Epilogue
movq %rbp,%rsp # Set %rsp to beginning of frame
popq %rbp      # Restore saved %rbp
ret
```



- The **function prologue is the process of creating a stack frame to hold callee function information.**
 - It is done by the callee function—the code to create the frame is located at the start of the callee function.
 - There are three steps to the function prologue:
 - 1 The current value of `%rbp` is pushed onto the stack. This will allow the calling function's stack frame to be rebuilt after the callee function finish;
 - 2 The current value of `%rsp` is moved into `%rbp`;
 - 3 Space is allocated for any local variables. This is done by subtracting their collective size (in hexadecimal form) from `%rsp`.
- Function epilogue **reverses the actions of the function prologue and returns control to the calling function by resetting its stack frame.**
 - The function epilogue also has three steps:
 - 1 `%rbp` is moved into `%rsp`;
 - 2 `%rbp` is popped from the stack;
 - 3 The return address is read from the top of the stack (where `%rbp` is pointing) and the instruction pointer jumps to that address.

Managing Local data

- Accessing to **stack frame data using %rsp** as anchor could be **dangerous**, because **there are instructions that change %rsp**.
- Using **Prologue** and **Epilogue** approach to manage stack frames, **provides a more stable option is to be used as acnchor: %rbp**

```
void fn(){  
    long x;  
    long y;  
    ...  
    x = 10;  
    y = 20;  
}
```

```
fn:  
    pushq %rbp  
    movq %rsp, %rbp  
    ...  
    subq $16, %rsp  
    ...  
    movq $20, -16(%rbp)  
    movq $10, -8(%rbp)  
  
    movq %rbp, %rsp  
    popq %rbp  
    ret
```

0x7FFFFFFFFFEF

0x7FFFFFFFFFEF

...	0x7FFFFFFFFFFF
...	0x7FFFFFFFFFF7F
	0x7FFFFFFFFFEF
	0x7FFFFFFFFFE7
	0x7FFFFFFFFDF
	0x7FFFFFFFFD7

%rbp0x7FFFFFFFFFE7

%rsp0x7FFFFFFFFD7

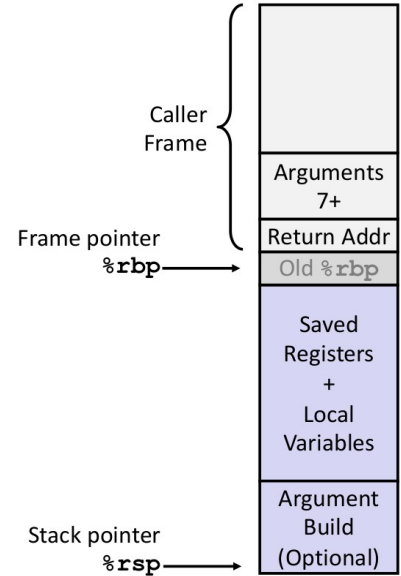
x

y

...	0x7FFFFFFFFFFF
...	0x7FFFFFFFFFF7F
0x7FFFFFFFFFFF	0x7FFFFFFFFFEF
10	0x7FFFFFFFFFE7
20	0x7FFFFFFFFDF
	0x7FFFFFFFFD7

Stack frame structure

- **Current stack frame** (Top to Bottom)
 - Argument build: 7+ parameters for function about to call
 - Local variables, if can't keep in registers
 - Saved register context
 - Old frame pointer
- **Caller stack frame**
 - Return address (pushed by `call`)
 - 7+ arguments for this call



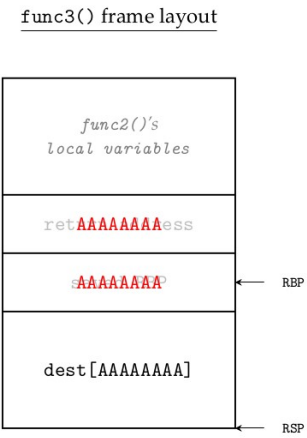
Memory Errors Exploitation

- A **buffer overflow** is the result of stuffing more data into a buffer than it can handle
- `strcpy`, is the most **infamous for being the cause of buffer overflows**.
- `strncpy` operates in the same way as `strcpy`, except that it copies a specified amount of bytes, `n`, from `src` to `dest`.
- Although the `strcpy` copy could be stopped before if the source string ends, under an attack the input length is controlled by attackers and will never happen.

```
char dest[8];

strcpy(dest, "AAAAAAAAAAAAAAAAAAAA");

strncpy(dest, "AAAAAAAAAAAAAAAAAAAA", 24);
```




- One of the first goals for attackers **when probing for buffer overflow vulnerabilities is gaining the ability to overwrite the stack frame return address.**
- When it is possible to **overwrite the return address of a stack frame**, and an attacker does so, the **CPU will jump to whatever address is stored in the return address when the function attempts to return to its caller**
- An example of `shellcode` that can be inserted into a vulnerable process to reboot a Linux x86-64 machine

```
char shellcode_reboot[] =  
"\xBA\xDC\xFE\x21\x43"  
"\xBE\x69\x19\x12\x28"  
"\xBF\xAD\xDE\xE1\xFE"  
"\xB0\xA9"  
"\x0F\x05";
```

Memory Protection Techniques: Stack Smashing Protector (SSP)

- To accomplish return address overwritten mitigation, a **canary value** was inserted next to the return address of the current stack frame to prevent an attacker from overwriting the return address.
- The **canary value is checked before the instruction pointer loads the return address of the stack frame.**
 - If the canary value is altered, the processor knows that an attack has been attempted and execution is aborted.



Memory Protection Techniques: Address Space Layout Randomisation (ASLR)

- ASLR is a protection technique that attempts **to render exploits that depend on predetermined memory addresses useless**
- It is a protection technique that which the memory address layout to prevent attacks that relies on knowing the location of an application's memory map