

Machine-Level Programming: Compound data types

Arquitectura de Computadores
Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto


Luís Nogueira (lmn@isep.ipp.pt)

Today

- **Structures**
 - Allocation
 - Access
- **Data alignment**
- **Unions**

Structure allocation

```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
};
```



■ Concept

- Contiguously-allocated region of memory
- Refer to members within structure by names
- Members may be of different types
- Compiler determines overall size and positions of fields

■ Fields ordered according to declaration

- Even if another ordering could yield a more compact representation

Declaring `struct` variables in C

■ `struct rec r1, r2, r3;`

- Declares and sets aside storage for three variables – `r1`, `r2`, and `r3` – each of type `struct rec`

■ `struct rec *ptr;`


- Declares a pointer to an object of type `struct rec`

■ `struct rec vec[25];`

- Declares a 25-element array of `struct rec`
- Allocates 25 units of storage, each one big enough to hold the data of one *struct rec*

Structure access in C


```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
};
```



- Given an instance of a structure: `struct rec r1;`
 - `r1.i = val;`
- Given a *pointer* to a structure: `struct rec* r = &r1;`
 - Using `*` and `.` operators: `(*r).i = val;`
 - Or, use `->` operator for short: `r->i = val;`

Structure access in Assembly


```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
};
```



- Machine-level program has no understanding of structures in source C code
- Address indicates first byte of structure
- Access elements with **offsets to the address** of the structure

Structure access in Assembly

```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
};
```



```
void set_i(struct rec *r, long val){  
    r->i = val;  
}
```


■ Access field with corresponding offset

```
# r in %rdi, val in %rsi  
movq %rsi, 16(%rdi)
```

```
# Mem[r+16] = val
```

Address of structure member

```
struct rec {  
    int a[4];  
    long i;  
    struct rec *next;  
};
```



```
int *get_ap(struct rec *r, int idx)  
{  
    return &(r->a[idx]);  
}
```

■ Add the field's offset to the structure's address

```
# r in %rdi, idx in %esi  
leaq    (%rdi,%rsi,4), %rax
```


Example: Following linked list

```


long length(struct rec*r){
    long len = 0;
    while (r) {
        len ++;
        r = r->next;
    }
    return len;
}

```

```

struct rec {
    int a[4];
    long i;
    struct rec *next;
};

```



```

# r in %rdi
.L11:                                # loop:
    addq    $1, %rax                # len ++
    movq    24(%rdi), %rdi          # r = Mem[r+24]
    testq   %rdi, %rdi              # Test r
    jne     .L11                    # If != 0, goto loop

```

Today

- Structures
 - Allocation
 - Access
- **Data alignment**
- Unions

Motivation

- Many computer systems place **restrictions on the allowable addresses** for the primitive data types
 - Require that the address for some type of object must be a multiple of some value K (typically 1, 2, 4, or 8)
- The x86-64 hardware will work correctly regardless of the alignment of data
- However, Intel recommends that data be aligned
 - Simplifies the design of the hardware forming the interface between the processor and the memory system
 - Improves memory system performance

Alignment principles

■ Aligned data

- Primitive data type requires K bytes
- Address must be multiple of K

■ Align enforcement by compiler

- Places `.align` directives in the assembly code indicating the desired alignment for global data
- Library routines that allocate memory, such as `malloc`, must be designed so that they return a pointer that satisfies the worst-case alignment restriction (typically 16 bytes on x86-84)
- For code involving structures, **inserts gaps in structure** to ensure correct alignment of fields

■ Treated differently by IA32 Linux, x86-64 Linux, Windows and Mac OS!

Specific cases of alignment (x86-64)

- **1 byte: char**
 - No restrictions on address
- **2 bytes: short**
 - Lowest 1 bit of address must be 0₂
- **4 bytes: int, float, ...**
 - Lowest 2 bits of address must be 00₂
- **8 bytes: long, double, char *, ...**
 - Lowest 3 bits of address must be 000₂

Data alignment in structures

■ Inside the structure


- The **alignment** requirement of **each field** type must be satisfied

■ Placement of structure in memory

- Each structure has an alignment requirement of K , where K is the size of the largest member of the structure
- The **starting address** of the structure must be a **multiple of K**
- The **total size** of the structure must be a **multiple of K**


Alignment requirements of members

■ Unaligned data



```
struct S1 {
    char c;
    int i[2];
    long v;
} *p;
```

■ Aligned data under x86-64



■ The compiler may need to add padding

- Between the fields (internal padding)
- At the end of the structure (external padding)

Alignment requirements of members

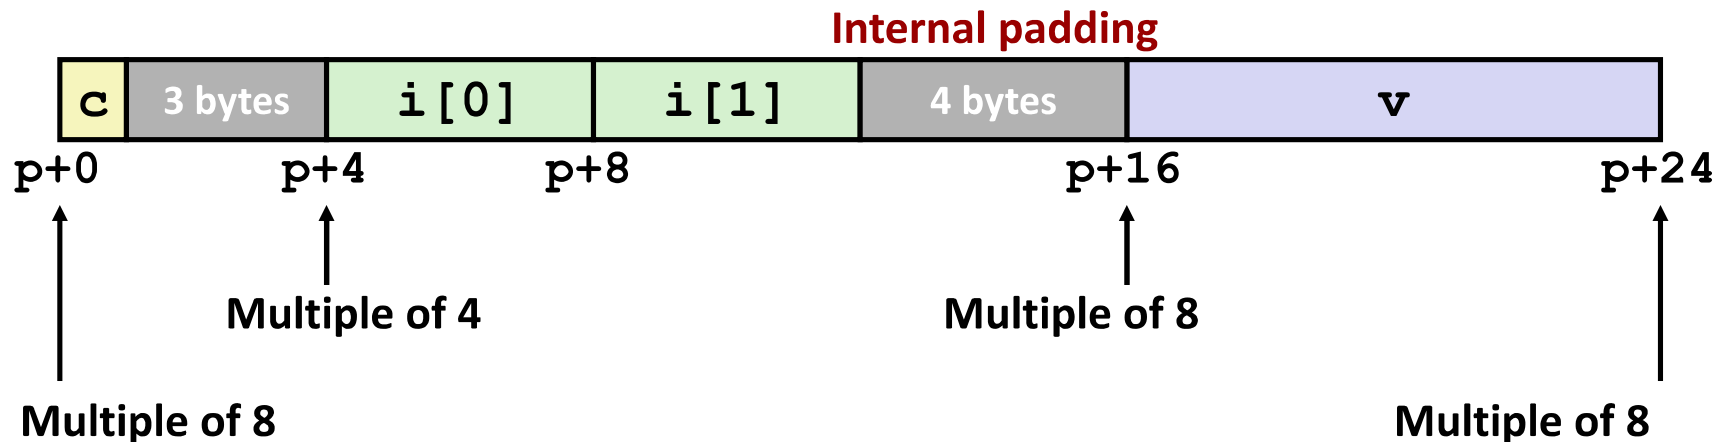
■ For each member

- Satisfy alignment requirement

```
struct S1 {
    char c;
    int i[2];
    long v;
} *p;
```

■ Overall structure placement

- Initial address and total size must be multiple of K
- K = 8, largest alignment of any member, due to `long` element

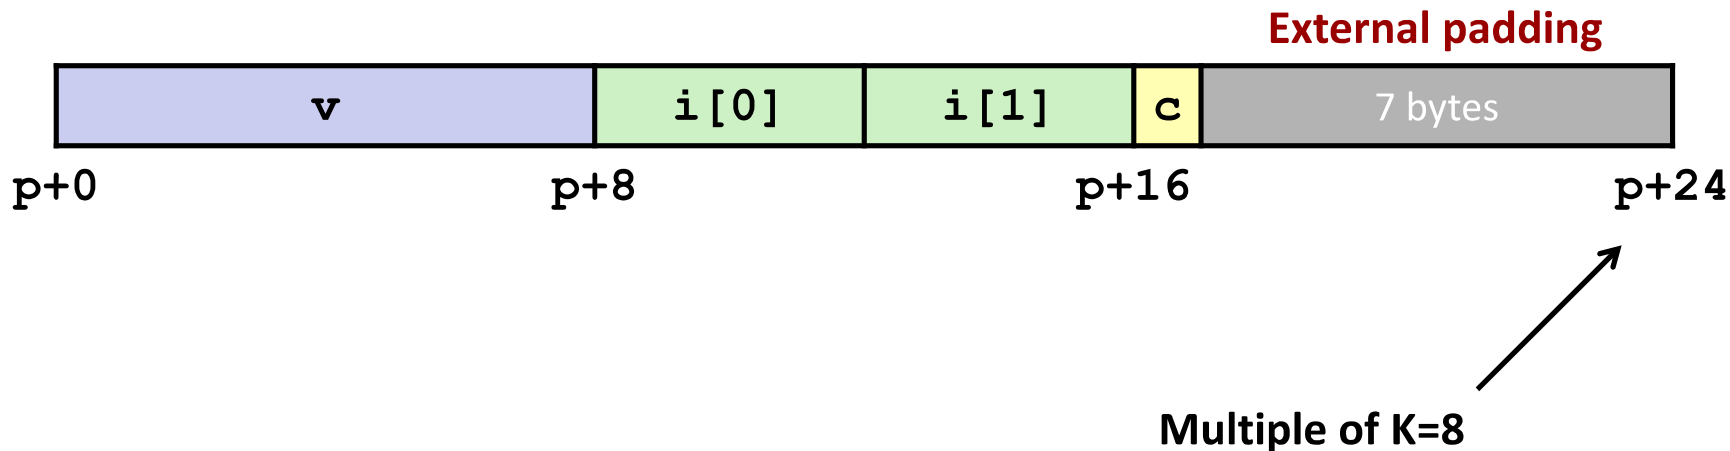


Overall alignment requirement

■ Overall structure placement

- Initial address and total size must be multiple of K
- K = 8, largest alignment of any member, due to long element

```
struct S2 {  
    long v;  
    int i[2];  
    char c;  
} *p;
```




Arrays of structures

```
struct S2 {
    long v;
    int i[2];
    char c;
} a[10];
```

■ Contiguously allocated region of $n * \text{sizeof}(\text{struct } x)$ bytes

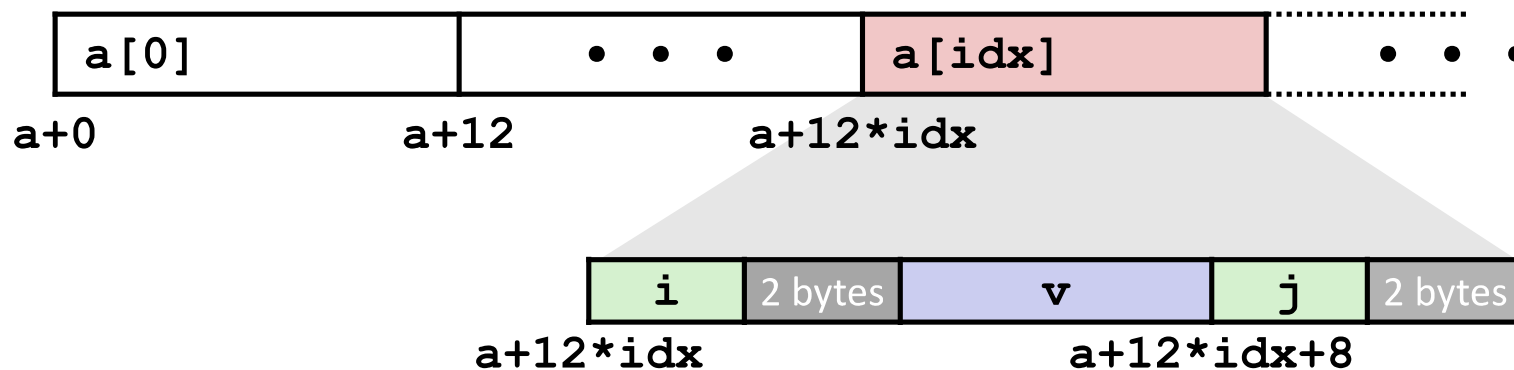
- Overall structure length multiple of K
- Satisfy alignment requirement for every element



Accessing array elements

- Compute array offset $12 * \text{idx}$
 - `sizeof(S3) = 12`, including alignment padding
- Element `j` is at offset 8 within structure

```
struct S3 {
    short i;
    int v;
    short j;
} a[10];
```



```
short get_j(struct s3 *a,
            int idx){
    return a[idx].j;
}
```

```
# 3*idx
leaq (%rsi,%rsi,2),%rax
# a+12*idx+8
movw 8(%rdi,%rax,4),%ax
```

Example exam question

Problem 5. (8 points):

Struct alignment. Consider the following C struct declaration:

```
typedef struct {
    char a;
    long b;
    float c;
    char d[3];
    int *e;
    short *f;
} foo;
```

1. Show how `foo` would be allocated in memory on an x86-64 Linux system. Label the bytes with the names of the various fields and **clearly mark the end of the struct**. Use an X to denote space that is allocated in the struct as padding.

```
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```


Example exam question

Problem 5. (8 points):

Struct alignment. Consider the following C struct declaration:

```
typedef struct {
    char a;
    long b;
    float c;
    char d[3];
    int *e;
    short *f;
} foo;
```


1. Show how `foo` would be allocated in memory on an x86-64 Linux system. Label the bytes with the names of the various fields and **clearly mark the end of the struct**. Use an `X` to denote space that is allocated in the struct as padding.



Saving space


- Put large data types first

```
struct S4 {
    char c;
    int i;
    char d;
} *p;
```




```
struct S5 {
    int i;
    char c;
    char d;
} *p;
```

- Effect on x86-64 (K=4)



12 bytes




8 bytes

- This strategy **can** save space in certain structures

Saving space



■ Put large data types first

```
struct S6 {
    char c;
    int i[2];
    long v;
} *p1;
```



```
struct S7 {
    long v;
    int i[2];
    char c;
} *p2;
```

■ Effect on x86-64 (K=8)




Unfortunately, doesn't satisfy requirement that struct's *total* size is a multiple of K

Saving space



■ Put large data types first

```
struct S6 {
    char c;
    int i[2];
    long v;
} *p1;
```



```
struct S7 {
    long v;
    int i[2];
    char c;
} *p2;
```

■ Effect on x86-64 (K=8)



Example exam question (Cont'd)

Problem 5. (8 points):

Struct alignment. Consider the following C struct declaration:

```
typedef struct {
    char a;
    long b;
    float c;
    char d[3];
    int *e;
    short *f;
} foo;
```

2. Rearrange the elements of `foo` to conserve the most space in memory. Label the bytes with the names of the various fields and **clearly mark the end of the struct**. Use an X to denote space that is allocated in the struct as padding.

```
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
|  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |  |
+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+--+
```

Example exam question (Cont'd)

Problem 5. (8 points):

Struct alignment. Consider the following C struct declaration:

```
typedef struct {
    char a;
    long b;
    float c;
    char d[3];
    int *e;
    short *f;
} foo;
```

2. Rearrange the elements of `foo` to conserve the most space in memory. Label the bytes with the names of the various fields and **clearly mark the end of the struct**. Use an X to denote space that is allocated in the struct as padding.



Today

- **Structures**
 - Allocation
 - Access
- **Data alignment**
- **Unions**

Unions



- **Provide a way to circumvent the type system of C**
 - Allow a single object to be referenced according to multiple types
- **The syntax of a union declaration is identical to that for structures, but its semantics are very different**
 - Rather than having the different fields reference different blocks of memory, they all reference the same block
- **The overall size of a union equals the maximum size of any of its fields**
 - Important when memory is scarcer, such as in embedded systems

Union allocation

- Allocated according to largest element
- Can only use one member at a time

```
union U1 {
    char c;
    int i[2];
    long v;
} *up;
```

```
struct S1 {
    char c;
    int i[2];
    long v;
} *sp;
```



Accessing union members

```
union Data{
    int i;
    float f;
    char str[20];
} data;

data.i = 10;
printf("data.i      : %d\n", data.i);      /* 10 */

data.f = 220.5;
printf("data.f      : %f\n", data.f);      /* 220.5 */

strcpy(data.str, "C Programming");
printf("data.str   : %s\n", data.str);     /* C Programming */
```

Accessing union members

```
union Data{
    int i;
    float f;
    char str[20];
} data;

data.i = 10;
data.f = 220.5;
strcpy(data.str, "C Programming");

printf("data.i      : %d\n", data.i);      /* 1917853763 */

printf("data.f      : %f\n", data.f);      /* 4122360580327.00 */

printf("data.str   : %s\n", data.str);     /* C Programming */
```

Accessing union members

- Unions can be useful in several contexts, but they can also lead to nasty bugs
- No checking is done to make sure that the right sort of use is made of the members
- It's up to the programmer to keep track of whatever type is put into it

Common use

■ Combining them with *structs*

- Produce variable types based on a common core, but then varying in parts where appropriate

```
typedef union {  
    int units;  
    float kgs;  
} amount;  
  
typedef struct {  
    char selling[15];  
    float unitprice;  
    int unittype;  
    amount howmuch;  
} product;
```

- Still require a lot of discipline to ensure that only the active member is accessed at any time

Summary

■ Structures

- Elements packed into single region of memory
- Access using offsets determined by compiler
- Possible require internal and external padding to ensure alignment

■ Combinations

- Can nest structure and array code arbitrarily

■ Unions

- Overlay declarations
- Way to circumvent type system